

```
In [2]: from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
import os
os.chdir("/content/gdrive/My Drive/Colab Notebooks/Startpkg_A2")
```

Mounted at /content/gdrive

```
In [3]: import numpy as np
import matplotlib.pyplot as plt

from utils.data_process import get_CIFAR10_data
from models.neural_net import NeuralNetwork
from kaggle_submission import output_submission_csv

##### If not using Colab, you may skip these setup #####
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
##### End of setup #####
```

## Loading CIFAR-10

Now that you have implemented a neural network that passes gradient checks and works on toy data, you will test your network on the CIFAR-10 dataset.

```
In [5]: # You can change these numbers for experimentation
# For submission be sure they are set to the default values
TRAIN_IMAGES = 49000
VAL_IMAGES = 1000
TEST_IMAGES = 5000 # Default is 5000, do not modify this for your submission.

data = get_CIFAR10_data(TRAIN_IMAGES, VAL_IMAGES, TEST_IMAGES)
X_train, y_train = data['X_train'], data['y_train']
X_val, y_val = data['X_val'], data['y_val']
X_test, y_test = data['X_test'], data['y_test']
```

# Train a network

To train our network we will use SGD. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

You can try different numbers of layers and also the different activation functions that you implemented on the CIFAR-10 dataset below.

```
In [6]: input_size = 32 * 32 * 3
num_layers = 3
hidden_size = 120
hidden_sizes = [hidden_size]*(num_layers-1)
num_classes = 10
learning_rate = 0.05
learning_rate_decay = 0.95
net = NeuralNetwork(input_size, hidden_sizes, num_classes, num_layers, nonline
arity='sigmoid')

# Train the network
stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=1000, batch_size=200,
                  learning_rate=learning_rate, learning_rate_decay=learning_rate_dec
ay,
                  reg=0.00, verbose=True)

# Predict on the validation set
val_acc = (net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)


iteration 0 / 1000: loss 2.360263
iteration 100 / 1000: loss 2.171381
iteration 200 / 1000: loss 2.022783
iteration 300 / 1000: loss 2.017632
iteration 400 / 1000: loss 1.983689
iteration 500 / 1000: loss 1.900510
iteration 600 / 1000: loss 1.885977
iteration 700 / 1000: loss 1.860290
iteration 800 / 1000: loss 1.836469
iteration 900 / 1000: loss 1.893836
Validation accuracy: 0.375
```

# Graph loss and train/val accuracies

Examining the loss graph along with the train and val accuracy graphs should help you gain some intuition for the hyperparameters you should try in the hyperparameter tuning below. It should also help with debugging any issues you might have with your network.

```
In [7]: # Plot the loss function and train / validation accuracies
plt.subplot(2, 1, 1)
plt.plot(stats['loss_history'])
plt.title('Loss history')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
plt.plot(stats['train_acc_history'], label='train')
plt.plot(stats['val_acc_history'], label='val')
plt.title('Classification accuracy history')
plt.xlabel('Epoch')
plt.ylabel('Classification accuracy')
plt.legend()
plt.show()
```



# Hyperparameter tuning

Once you have successfully trained a network you can tune your hyperparameters to increase your accuracy.

Based on the graphs of the loss function above you should be able to develop some intuition about what hyperparameter adjustments may be necessary. A very noisy loss implies that the learning rate might be too high, while a linearly decreasing loss would suggest that the learning rate may be too low. A large gap between training and validation accuracy would suggest overfitting due to large model without much regularization. No gap between training and validation accuracy would indicate low model capacity (low model complexity).

You will compare networks of two and three layers using the different activation functions you implemented.

The different hyperparameters you can experiment with are:

- **Batch size:** We recommend you leave this at 200 initially which is the batch size we used.
- **Number of iterations:** You can gain an intuition for how many iterations to run by checking when the validation accuracy plateaus in your train/val accuracy graph.
- **Initialization** Weight initialization is very important for neural networks. We used the initialization  $w = \text{np.random.randn}(n) / \sqrt{n}$  where  $n$  is the input dimension for layer corresponding to  $w$ . We recommend you stick with the given initializations, but you may explore modifying these. Typical initialization practices: <http://cs231n.github.io/neural-networks-2/#init> (<http://cs231n.github.io/neural-networks-2/#init>)
- **Learning rate:** Generally from around  $1e-4$  to  $1e-1$  is a good range to explore according to our implementation.
- **Learning rate decay:** We recommend a 0.95 decay to start.
- **Hidden layer size:** You should explore up to around 120 units per layer. For three-layer network, we fixed the two hidden layers to be the same size when obtaining the target numbers. However, you may experiment with having different size hidden layers.
- **Regularization coefficient:** We recommend trying values in the range 0 to 0.1.

Hints:

- After getting a sense of the parameters by trying a few values yourself, you will likely want to write a few for loops to traverse over a set of hyperparameters.
- If you find that your train loss is decreasing, but your train and val accuracy start to decrease rather than increase, your model likely started minimizing the regularization term. To prevent this you will need to decrease the regularization coefficient.

## DEFINING MATRIX OF PARAMETERS (ADDED BY OR)

```

In [8]: #####
#####
###
#                               DEFINING THE MATRIX OF PARAMETERS
#
#                               THIS MATRIX WILL BE USED FOR OPTIMIZATION OF ALL THE MODELS
#
#####
#####

learning_rate_values = [0.1, 0.05, 0.03]
hidden_size_values = [20, 50, 80, 100, 120]
batch_size_values = [50, 100, 200]

#All of this is to build the parameter matrix

#Initializing the parameters matrix

parameters_matrix = np.zeros((len(learning_rate_values)*len(hidden_size_values)
)*len(batch_size_values),3))

#Entering the learning rate values

parameters_matrix[0:len(hidden_size_values)*len(batch_size_values)][:,0] = lea
rning_rate_values[0]
parameters_matrix[len(hidden_size_values)*len(batch_size_values):2*len(hidden_
size_values)*len(batch_size_values)][:,0] = learning_rate_values[1]
parameters_matrix[2*len(hidden_size_values)*len(batch_size_values):3*len(hidde
n_size_values)*len(batch_size_values)][:,0] = learning_rate_values[2]

#Entering the hidden size values

parameters_matrix[0:len(parameters_matrix):5,1] = hidden_size_values[0]
parameters_matrix[1:len(parameters_matrix):5,1] = hidden_size_values[1]
parameters_matrix[2:len(parameters_matrix):5,1] = hidden_size_values[2]
parameters_matrix[3:len(parameters_matrix):5,1] = hidden_size_values[3]
parameters_matrix[4:len(parameters_matrix):5,1] = hidden_size_values[4]

#Entering the hidden size values

parameters_matrix[0:len(parameters_matrix):3,2] = batch_size_values[0]
parameters_matrix[1:len(parameters_matrix):3,2] = batch_size_values[1]
parameters_matrix[2:len(parameters_matrix):3,2] = batch_size_values[2]

```

## Two-layer Relu Activation Network

In [9]: `best_2layer_relu = None` # store the best model into this

```
#####  
###  
# TODO: Tune hyperparameters using the validation set. Store your best trained  
#  
# model in best_2layer_relu.  
#  
#####  
###  
  
#Initialize "best" validation accuracy  
  
best_val_acc = 0  
  
#Initialize vector to store all validation accuracies  
  
val_acc_vec = np.zeros(len(parameters_matrix))  
  
#Cycle through the matrix of parameters  
  
for i in range(len(parameters_matrix)):  
  
    input_size = 32 * 32 * 3  
    num_layers = 2  
    hidden_size = int(parameters_matrix[i,1])  
    hidden_sizes = [hidden_size]*(num_layers-1)  
    num_classes = 10  
    learning_rate = parameters_matrix[i,0]  
    learning_rate_decay = 0.95  
    batch_size = int(parameters_matrix[i,2])  
    net = NeuralNetwork(input_size, hidden_sizes, num_classes, num_layers, nonli  
nearity='relu')  
  
    # Train the network  
    stats = net.train(X_train, y_train, X_val, y_val,  
                      num_iters=1000, batch_size=batch_size,  
                      learning_rate=learning_rate, learning_rate_decay=learning_rate_d  
ecay,  
                      reg=0.0, verbose=True)  
  
    # Predict on the validation set  
    val_acc = (net.predict(X_val) == y_val).mean()  
  
    print('For parameters: ', parameters_matrix[i])  
    print('Validation accuracy: ', val_acc)  
  
    #Store the validation accuracies  
  
    val_acc_vec[i] = val_acc  
  
    #Save the best model and stats  
  
    if val_acc > best_val_acc:  
  
        best_val_acc = val_acc
```

```
best_2layer_relu = net  
best_stats = stats
```

```
#Print the best combination of trial parameters:
```

```
ind = np.argmax(val_acc_vec)
```

```
print('The best combination of parameters is: learning rate = ',parameters_mat  
rix[ind,0], ', hidden layer size = ',int(parameters_matrix[ind,1]), ', batch s  
ize = ',int(parameters_matrix[ind,2]))  
print('For this combination, the validation accuracy is ', val_acc_vec[ind])
```

iteration 0 / 1000: loss 2.715286  
iteration 100 / 1000: loss 1.866864  
iteration 200 / 1000: loss 1.913203  
iteration 300 / 1000: loss 1.814135  
iteration 400 / 1000: loss 1.981856  
iteration 500 / 1000: loss 1.978752  
iteration 600 / 1000: loss 1.811956  
iteration 700 / 1000: loss 1.678494  
iteration 800 / 1000: loss 1.764641  
iteration 900 / 1000: loss 1.839433  
For parameters: [ 0.1 20. 50. ]  
Validation accuracy: 0.398  
iteration 0 / 1000: loss 2.542021  
iteration 100 / 1000: loss 1.683022  
iteration 200 / 1000: loss 1.664625  
iteration 300 / 1000: loss 1.557218  
iteration 400 / 1000: loss 1.844410  
iteration 500 / 1000: loss 1.750250  
iteration 600 / 1000: loss 1.544880  
iteration 700 / 1000: loss 1.627553  
iteration 800 / 1000: loss 1.524133  
iteration 900 / 1000: loss 1.440771  
For parameters: [ 0.1 50. 100. ]  
Validation accuracy: 0.473  
iteration 0 / 1000: loss 2.428241  
iteration 100 / 1000: loss 1.796744  
iteration 200 / 1000: loss 1.546376  
iteration 300 / 1000: loss 1.577306  
iteration 400 / 1000: loss 1.571245  
iteration 500 / 1000: loss 1.298537  
iteration 600 / 1000: loss 1.566067  
iteration 700 / 1000: loss 1.336600  
iteration 800 / 1000: loss 1.361602  
iteration 900 / 1000: loss 1.421368  
For parameters: [1.e-01 8.e+01 2.e+02]  
Validation accuracy: 0.489  
iteration 0 / 1000: loss 2.434387  
iteration 100 / 1000: loss 1.897614  
iteration 200 / 1000: loss 1.790517  
iteration 300 / 1000: loss 1.549533  
iteration 400 / 1000: loss 1.756254  
iteration 500 / 1000: loss 1.785943  
iteration 600 / 1000: loss 1.959229  
iteration 700 / 1000: loss 1.622413  
iteration 800 / 1000: loss 1.623365  
iteration 900 / 1000: loss 2.029529  
For parameters: [ 0.1 100. 50. ]  
Validation accuracy: 0.414  
iteration 0 / 1000: loss 2.376951  
iteration 100 / 1000: loss 1.788000  
iteration 200 / 1000: loss 1.559161  
iteration 300 / 1000: loss 1.560658  
iteration 400 / 1000: loss 1.558631  
iteration 500 / 1000: loss 1.715426  
iteration 600 / 1000: loss 1.543630  
iteration 700 / 1000: loss 1.522090  
iteration 800 / 1000: loss 1.422984



iteration 900 / 1000: loss 1.484674  
For parameters: [1.0e-01 1.2e+02 1.0e+02]  
Validation accuracy: 0.468  
iteration 0 / 1000: loss 2.426119  
iteration 100 / 1000: loss 1.717044  
iteration 200 / 1000: loss 1.629928  
iteration 300 / 1000: loss 1.697360  
iteration 400 / 1000: loss 1.646700  
iteration 500 / 1000: loss 1.741680  
iteration 600 / 1000: loss 1.713300  
iteration 700 / 1000: loss 1.547610  
iteration 800 / 1000: loss 1.525743  
iteration 900 / 1000: loss 1.465556  
For parameters: [1.e-01 2.e+01 2.e+02]  
Validation accuracy: 0.441  
iteration 0 / 1000: loss 2.662212  
iteration 100 / 1000: loss 1.804719  
iteration 200 / 1000: loss 1.825623  
iteration 300 / 1000: loss 1.827037  
iteration 400 / 1000: loss 1.991980  
iteration 500 / 1000: loss 1.731884  
iteration 600 / 1000: loss 1.457862  
iteration 700 / 1000: loss 1.559146  
iteration 800 / 1000: loss 1.554897  
iteration 900 / 1000: loss 1.445868  
For parameters: [ 0.1 50. 50. ]  
Validation accuracy: 0.36  
iteration 0 / 1000: loss 2.518125  
iteration 100 / 1000: loss 1.738500  
iteration 200 / 1000: loss 1.639498  
iteration 300 / 1000: loss 1.994568  
iteration 400 / 1000: loss 1.681102  
iteration 500 / 1000: loss 1.672262  
iteration 600 / 1000: loss 1.432016  
iteration 700 / 1000: loss 1.425032  
iteration 800 / 1000: loss 1.553365  
iteration 900 / 1000: loss 1.344180  
For parameters: [ 0.1 80. 100. ]  
Validation accuracy: 0.418  
iteration 0 / 1000: loss 2.380153  
iteration 100 / 1000: loss 1.728520  
iteration 200 / 1000: loss 1.534835  
iteration 300 / 1000: loss 1.449737  
iteration 400 / 1000: loss 1.362949  
iteration 500 / 1000: loss 1.485556  
iteration 600 / 1000: loss 1.462125  
iteration 700 / 1000: loss 1.338245  
iteration 800 / 1000: loss 1.258395  
iteration 900 / 1000: loss 1.283329  
For parameters: [1.e-01 1.e+02 2.e+02]  
Validation accuracy: 0.5  
iteration 0 / 1000: loss 2.525483  
iteration 100 / 1000: loss 1.878190  
iteration 200 / 1000: loss 1.749712  
iteration 300 / 1000: loss 1.666091  
iteration 400 / 1000: loss 1.970726  
iteration 500 / 1000: loss 2.062984

iteration 600 / 1000: loss 1.614487  
iteration 700 / 1000: loss 1.638533  
iteration 800 / 1000: loss 1.766345  
iteration 900 / 1000: loss 1.767084  
For parameters: [1.0e-01 1.2e+02 5.0e+01]  
Validation accuracy: 0.418  
iteration 0 / 1000: loss 2.606611  
iteration 100 / 1000: loss 1.730961  
iteration 200 / 1000: loss 1.743571  
iteration 300 / 1000: loss 1.783963  
iteration 400 / 1000: loss 1.847133  
iteration 500 / 1000: loss 1.416942  
iteration 600 / 1000: loss 1.683316  
iteration 700 / 1000: loss 1.667880  
iteration 800 / 1000: loss 1.747269  
iteration 900 / 1000: loss 1.620577  
For parameters: [ 0.1 20. 100. ]  
Validation accuracy: 0.43  
iteration 0 / 1000: loss 2.498759  
iteration 100 / 1000: loss 1.711566  
iteration 200 / 1000: loss 1.572389  
iteration 300 / 1000: loss 1.591275  
iteration 400 / 1000: loss 1.555117  
iteration 500 / 1000: loss 1.541600  
iteration 600 / 1000: loss 1.518847  
iteration 700 / 1000: loss 1.381872  
iteration 800 / 1000: loss 1.422076  
iteration 900 / 1000: loss 1.315779  
For parameters: [1.e-01 5.e+01 2.e+02]  
Validation accuracy: 0.477  
iteration 0 / 1000: loss 2.480604  
iteration 100 / 1000: loss 1.962881  
iteration 200 / 1000: loss 1.836400  
iteration 300 / 1000: loss 1.554923  
iteration 400 / 1000: loss 1.772844  
iteration 500 / 1000: loss 1.915968  
iteration 600 / 1000: loss 1.679191  
iteration 700 / 1000: loss 1.792873  
iteration 800 / 1000: loss 1.380780  
iteration 900 / 1000: loss 1.774523  
For parameters: [ 0.1 80. 50. ]  
Validation accuracy: 0.44  
iteration 0 / 1000: loss 2.475962  
iteration 100 / 1000: loss 1.697423  
iteration 200 / 1000: loss 1.684923  
iteration 300 / 1000: loss 1.464603  
iteration 400 / 1000: loss 1.566533  
iteration 500 / 1000: loss 1.393598  
iteration 600 / 1000: loss 1.483803  
iteration 700 / 1000: loss 1.518905  
iteration 800 / 1000: loss 1.502741  
iteration 900 / 1000: loss 1.444068  
For parameters: [ 0.1 100. 100. ]  
Validation accuracy: 0.466  
iteration 0 / 1000: loss 2.507818  
iteration 100 / 1000: loss 1.661972  
iteration 200 / 1000: loss 1.693131

iteration 300 / 1000: loss 1.489109  
iteration 400 / 1000: loss 1.571926  
iteration 500 / 1000: loss 1.349370  
iteration 600 / 1000: loss 1.234844  
iteration 700 / 1000: loss 1.226016  
iteration 800 / 1000: loss 1.545961  
iteration 900 / 1000: loss 1.250473  
For parameters: [1.0e-01 1.2e+02 2.0e+02]  
Validation accuracy: 0.524  
iteration 0 / 1000: loss 2.649751  
iteration 100 / 1000: loss 2.082025  
iteration 200 / 1000: loss 1.839678  
iteration 300 / 1000: loss 1.873448  
iteration 400 / 1000: loss 1.777536  
iteration 500 / 1000: loss 1.778407  
iteration 600 / 1000: loss 1.748325  
iteration 700 / 1000: loss 1.654586  
iteration 800 / 1000: loss 1.835584  
iteration 900 / 1000: loss 1.731624  
For parameters: [ 0.05 20. 50. ]  
Validation accuracy: 0.399  
iteration 0 / 1000: loss 2.438073  
iteration 100 / 1000: loss 1.774453  
iteration 200 / 1000: loss 1.587868  
iteration 300 / 1000: loss 1.613218  
iteration 400 / 1000: loss 1.685128  
iteration 500 / 1000: loss 1.481976  
iteration 600 / 1000: loss 1.693441  
iteration 700 / 1000: loss 1.522465  
iteration 800 / 1000: loss 1.260546  
iteration 900 / 1000: loss 1.456370  
For parameters: [5.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.473  
iteration 0 / 1000: loss 2.609454  
iteration 100 / 1000: loss 1.655263  
iteration 200 / 1000: loss 1.566539  
iteration 300 / 1000: loss 1.617915  
iteration 400 / 1000: loss 1.508518  
iteration 500 / 1000: loss 1.480823  
iteration 600 / 1000: loss 1.466878  
iteration 700 / 1000: loss 1.452852  
iteration 800 / 1000: loss 1.281365  
iteration 900 / 1000: loss 1.332153  
For parameters: [5.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.498  
iteration 0 / 1000: loss 2.340921  
iteration 100 / 1000: loss 1.618954  
iteration 200 / 1000: loss 1.785193  
iteration 300 / 1000: loss 1.815763  
iteration 400 / 1000: loss 1.767731  
iteration 500 / 1000: loss 1.689865  
iteration 600 / 1000: loss 1.502381  
iteration 700 / 1000: loss 1.792459  
iteration 800 / 1000: loss 1.373051  
iteration 900 / 1000: loss 1.661950  
For parameters: [5.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.458

iteration 0 / 1000: loss 2.420423  
iteration 100 / 1000: loss 1.847190  
iteration 200 / 1000: loss 1.671877  
iteration 300 / 1000: loss 1.777128  
iteration 400 / 1000: loss 1.563582  
iteration 500 / 1000: loss 1.404980  
iteration 600 / 1000: loss 1.405071  
iteration 700 / 1000: loss 1.259767  
iteration 800 / 1000: loss 1.368514  
iteration 900 / 1000: loss 1.461776  
For parameters: [5.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.483  
iteration 0 / 1000: loss 2.461378  
iteration 100 / 1000: loss 1.787870  
iteration 200 / 1000: loss 1.709388  
iteration 300 / 1000: loss 1.716957  
iteration 400 / 1000: loss 1.573373  
iteration 500 / 1000: loss 1.588788  
iteration 600 / 1000: loss 1.542064  
iteration 700 / 1000: loss 1.532098  
iteration 800 / 1000: loss 1.496796  
iteration 900 / 1000: loss 1.462122  
For parameters: [5.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.445  
iteration 0 / 1000: loss 2.407305  
iteration 100 / 1000: loss 1.643106  
iteration 200 / 1000: loss 2.024314  
iteration 300 / 1000: loss 1.613249  
iteration 400 / 1000: loss 1.446963  
iteration 500 / 1000: loss 1.643224  
iteration 600 / 1000: loss 1.396599  
iteration 700 / 1000: loss 1.382745  
iteration 800 / 1000: loss 1.682422  
iteration 900 / 1000: loss 1.584512  
For parameters: [ 0.05 50. 50. ]  
Validation accuracy: 0.443  
iteration 0 / 1000: loss 2.412643  
iteration 100 / 1000: loss 1.810660  
iteration 200 / 1000: loss 1.513533  
iteration 300 / 1000: loss 1.638817  
iteration 400 / 1000: loss 1.532174  
iteration 500 / 1000: loss 1.508489  
iteration 600 / 1000: loss 1.628086  
iteration 700 / 1000: loss 1.277978  
iteration 800 / 1000: loss 1.360869  
iteration 900 / 1000: loss 1.420019  
For parameters: [5.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.469  
iteration 0 / 1000: loss 2.624492  
iteration 100 / 1000: loss 1.628949  
iteration 200 / 1000: loss 1.541047  
iteration 300 / 1000: loss 1.472769  
iteration 400 / 1000: loss 1.457646  
iteration 500 / 1000: loss 1.513345  
iteration 600 / 1000: loss 1.428303  
iteration 700 / 1000: loss 1.321445  
iteration 800 / 1000: loss 1.296052

iteration 900 / 1000: loss 1.409009  
For parameters: [5.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.5  
iteration 0 / 1000: loss 2.735369  
iteration 100 / 1000: loss 1.821706  
iteration 200 / 1000: loss 1.535711  
iteration 300 / 1000: loss 1.557933  
iteration 400 / 1000: loss 1.609216  
iteration 500 / 1000: loss 1.458676  
iteration 600 / 1000: loss 1.513089  
iteration 700 / 1000: loss 1.433682  
iteration 800 / 1000: loss 1.759217  
iteration 900 / 1000: loss 1.533537  
For parameters: [5.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.448  
iteration 0 / 1000: loss 2.558711  
iteration 100 / 1000: loss 1.891120  
iteration 200 / 1000: loss 1.714737  
iteration 300 / 1000: loss 1.686663  
iteration 400 / 1000: loss 1.608266  
iteration 500 / 1000: loss 1.538979  
iteration 600 / 1000: loss 1.469570  
iteration 700 / 1000: loss 1.824436  
iteration 800 / 1000: loss 1.551256  
iteration 900 / 1000: loss 1.894040  
For parameters: [5.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.43  
iteration 0 / 1000: loss 2.450733  
iteration 100 / 1000: loss 1.867561  
iteration 200 / 1000: loss 1.568302  
iteration 300 / 1000: loss 1.579026  
iteration 400 / 1000: loss 1.488592  
iteration 500 / 1000: loss 1.438244  
iteration 600 / 1000: loss 1.514497  
iteration 700 / 1000: loss 1.348599  
iteration 800 / 1000: loss 1.471637  
iteration 900 / 1000: loss 1.403929  
For parameters: [5.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.483  
iteration 0 / 1000: loss 2.804274  
iteration 100 / 1000: loss 1.446387  
iteration 200 / 1000: loss 1.714435  
iteration 300 / 1000: loss 1.557163  
iteration 400 / 1000: loss 1.360716  
iteration 500 / 1000: loss 1.639995  
iteration 600 / 1000: loss 1.620307  
iteration 700 / 1000: loss 1.712601  
iteration 800 / 1000: loss 1.598090  
iteration 900 / 1000: loss 1.427387  
For parameters: [5.e-02 8.e+01 5.e+01]  
Validation accuracy: 0.451  
iteration 0 / 1000: loss 2.474303  
iteration 100 / 1000: loss 1.678994  
iteration 200 / 1000: loss 1.629488  
iteration 300 / 1000: loss 1.396176  
iteration 400 / 1000: loss 1.691628  
iteration 500 / 1000: loss 1.627583

iteration 600 / 1000: loss 1.309060  
iteration 700 / 1000: loss 1.322883  
iteration 800 / 1000: loss 1.460409  
iteration 900 / 1000: loss 1.629806  
For parameters: [5.e-02 1.e+02 1.e+02]  
Validation accuracy: 0.472  
iteration 0 / 1000: loss 2.661310  
iteration 100 / 1000: loss 1.739933  
iteration 200 / 1000: loss 1.568916  
iteration 300 / 1000: loss 1.407349  
iteration 400 / 1000: loss 1.550511  
iteration 500 / 1000: loss 1.470718  
iteration 600 / 1000: loss 1.301667  
iteration 700 / 1000: loss 1.432915  
iteration 800 / 1000: loss 1.474231  
iteration 900 / 1000: loss 1.419195  
For parameters: [5.0e-02 1.2e+02 2.0e+02]  
Validation accuracy: 0.507  
iteration 0 / 1000: loss 2.478321  
iteration 100 / 1000: loss 1.991663  
iteration 200 / 1000: loss 1.918174  
iteration 300 / 1000: loss 1.721801  
iteration 400 / 1000: loss 1.792469  
iteration 500 / 1000: loss 1.754125  
iteration 600 / 1000: loss 2.038703  
iteration 700 / 1000: loss 1.731625  
iteration 800 / 1000: loss 1.648066  
iteration 900 / 1000: loss 1.600500  
For parameters: [3.e-02 2.e+01 5.e+01]  
Validation accuracy: 0.411  
iteration 0 / 1000: loss 2.555676  
iteration 100 / 1000: loss 1.851040  
iteration 200 / 1000: loss 1.783851  
iteration 300 / 1000: loss 1.622707  
iteration 400 / 1000: loss 1.575700  
iteration 500 / 1000: loss 1.636633  
iteration 600 / 1000: loss 1.540888  
iteration 700 / 1000: loss 1.532103  
iteration 800 / 1000: loss 1.642933  
iteration 900 / 1000: loss 1.487687  
For parameters: [3.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.479  
iteration 0 / 1000: loss 2.516289  
iteration 100 / 1000: loss 1.631177  
iteration 200 / 1000: loss 1.585127  
iteration 300 / 1000: loss 1.475524  
iteration 400 / 1000: loss 1.555140  
iteration 500 / 1000: loss 1.549211  
iteration 600 / 1000: loss 1.479004  
iteration 700 / 1000: loss 1.484843  
iteration 800 / 1000: loss 1.455360  
iteration 900 / 1000: loss 1.289475  
For parameters: [3.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.491  
iteration 0 / 1000: loss 2.526589  
iteration 100 / 1000: loss 1.678577  
iteration 200 / 1000: loss 1.616626

iteration 300 / 1000: loss 1.706297  
iteration 400 / 1000: loss 1.816006  
iteration 500 / 1000: loss 1.996626  
iteration 600 / 1000: loss 1.174053  
iteration 700 / 1000: loss 1.475161  
iteration 800 / 1000: loss 1.351806  
iteration 900 / 1000: loss 1.625959  
For parameters: [3.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.464  
iteration 0 / 1000: loss 2.451267  
iteration 100 / 1000: loss 1.872902  
iteration 200 / 1000: loss 1.718976  
iteration 300 / 1000: loss 1.510867  
iteration 400 / 1000: loss 1.567572  
iteration 500 / 1000: loss 1.536055  
iteration 600 / 1000: loss 1.454384  
iteration 700 / 1000: loss 1.620166  
iteration 800 / 1000: loss 1.520594  
iteration 900 / 1000: loss 1.520137  
For parameters: [3.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.475  
iteration 0 / 1000: loss 2.491889  
iteration 100 / 1000: loss 1.755961  
iteration 200 / 1000: loss 1.769272  
iteration 300 / 1000: loss 1.588391  
iteration 400 / 1000: loss 1.688738  
iteration 500 / 1000: loss 1.702200  
iteration 600 / 1000: loss 1.547895  
iteration 700 / 1000: loss 1.578232  
iteration 800 / 1000: loss 1.568111  
iteration 900 / 1000: loss 1.724036  
For parameters: [3.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.416  
iteration 0 / 1000: loss 2.665119  
iteration 100 / 1000: loss 2.003268  
iteration 200 / 1000: loss 1.749592  
iteration 300 / 1000: loss 1.840942  
iteration 400 / 1000: loss 2.000071  
iteration 500 / 1000: loss 1.545833  
iteration 600 / 1000: loss 1.693547  
iteration 700 / 1000: loss 1.438741  
iteration 800 / 1000: loss 1.539980  
iteration 900 / 1000: loss 1.552777  
For parameters: [3.e-02 5.e+01 5.e+01]  
Validation accuracy: 0.433  
iteration 0 / 1000: loss 2.631301  
iteration 100 / 1000: loss 1.838058  
iteration 200 / 1000: loss 1.506151  
iteration 300 / 1000: loss 1.747571  
iteration 400 / 1000: loss 1.568485  
iteration 500 / 1000: loss 1.609352  
iteration 600 / 1000: loss 1.480039  
iteration 700 / 1000: loss 1.310527  
iteration 800 / 1000: loss 1.458614  
iteration 900 / 1000: loss 1.369625  
For parameters: [3.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.472

iteration 0 / 1000: loss 2.661046  
iteration 100 / 1000: loss 1.756782  
iteration 200 / 1000: loss 1.621254  
iteration 300 / 1000: loss 1.510133  
iteration 400 / 1000: loss 1.586507  
iteration 500 / 1000: loss 1.560481  
iteration 600 / 1000: loss 1.540028  
iteration 700 / 1000: loss 1.459255  
iteration 800 / 1000: loss 1.304648  
iteration 900 / 1000: loss 1.386875  
For parameters: [3.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.491  
iteration 0 / 1000: loss 2.574094  
iteration 100 / 1000: loss 1.983784  
iteration 200 / 1000: loss 1.645728  
iteration 300 / 1000: loss 1.448059  
iteration 400 / 1000: loss 1.729623  
iteration 500 / 1000: loss 1.657289  
iteration 600 / 1000: loss 1.760682  
iteration 700 / 1000: loss 1.624786  
iteration 800 / 1000: loss 1.543331  
iteration 900 / 1000: loss 1.804036  
For parameters: [3.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.451  
iteration 0 / 1000: loss 2.360071  
iteration 100 / 1000: loss 2.012974  
iteration 200 / 1000: loss 1.779618  
iteration 300 / 1000: loss 1.784505  
iteration 400 / 1000: loss 1.615778  
iteration 500 / 1000: loss 1.592882  
iteration 600 / 1000: loss 1.653272  
iteration 700 / 1000: loss 1.398708  
iteration 800 / 1000: loss 1.820599  
iteration 900 / 1000: loss 1.541349  
For parameters: [3.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.426  
iteration 0 / 1000: loss 2.670196  
iteration 100 / 1000: loss 1.696313  
iteration 200 / 1000: loss 1.694602  
iteration 300 / 1000: loss 1.589026  
iteration 400 / 1000: loss 1.552249  
iteration 500 / 1000: loss 1.499694  
iteration 600 / 1000: loss 1.536549  
iteration 700 / 1000: loss 1.542548  
iteration 800 / 1000: loss 1.389237  
iteration 900 / 1000: loss 1.401183  
For parameters: [3.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.485  
iteration 0 / 1000: loss 2.539125  
iteration 100 / 1000: loss 2.057166  
iteration 200 / 1000: loss 1.845248  
iteration 300 / 1000: loss 1.528448  
iteration 400 / 1000: loss 1.707525  
iteration 500 / 1000: loss 1.816928  
iteration 600 / 1000: loss 1.528792  
iteration 700 / 1000: loss 1.604526  
iteration 800 / 1000: loss 1.640242



```
iteration 900 / 1000: loss 1.621079
For parameters: [3.e-02 8.e+01 5.e+01]
Validation accuracy: 0.455
iteration 0 / 1000: loss 2.432269
iteration 100 / 1000: loss 1.923649
iteration 200 / 1000: loss 1.557253
iteration 300 / 1000: loss 1.481281
iteration 400 / 1000: loss 1.744673
iteration 500 / 1000: loss 1.654912
iteration 600 / 1000: loss 1.789358
iteration 700 / 1000: loss 1.453594
iteration 800 / 1000: loss 1.354064
iteration 900 / 1000: loss 1.573829
For parameters: [3.e-02 1.e+02 1.e+02]
Validation accuracy: 0.464
iteration 0 / 1000: loss 2.622189
iteration 100 / 1000: loss 1.698678
iteration 200 / 1000: loss 1.612404
iteration 300 / 1000: loss 1.483058
iteration 400 / 1000: loss 1.544175
iteration 500 / 1000: loss 1.466680
iteration 600 / 1000: loss 1.372916
iteration 700 / 1000: loss 1.443018
iteration 800 / 1000: loss 1.327923
iteration 900 / 1000: loss 1.355483
For parameters: [3.0e-02 1.2e+02 2.0e+02]
Validation accuracy: 0.493
The best combination of parameters is: learning rate = 0.1 , hidden layer si
ze = 120 , batch size = 200
For this combination, the validation accuracy is 0.524
```

## Two-layer Sigmoid Activation Network

In [13]: `best_2layer_sigmoid = None` # store the best model into this

```
#####  
###  
# TODO: Tune hyperparameters using the validation set. Store your best trained  
#  
# model in best_2layer_sigmoid.  
#  
#####  
###  
  
#Initialize "best" validation accuracy  
  
best_val_acc = 0  
  
#Initialize vector to store all validation accuracies  
  
val_acc_vec = np.zeros(len(parameters_matrix))  
  
#Cycle through the matrix of parameters  
  
for i in range(len(parameters_matrix)):  
  
    input_size = 32 * 32 * 3  
    num_layers = 2  
    hidden_size = int(parameters_matrix[i,1])  
    hidden_sizes = [hidden_size]*(num_layers-1)  
    num_classes = 10  
    learning_rate = parameters_matrix[i,0]  
    learning_rate_decay = 0.95  
    batch_size = int(parameters_matrix[i,2])  
    net = NeuralNetwork(input_size, hidden_sizes, num_classes, num_layers, nonli  
nearity='sigmoid')  
  
    # Train the network  
    stats = net.train(X_train, y_train, X_val, y_val,  
                      num_iters=1000, batch_size=batch_size,  
                      learning_rate=learning_rate, learning_rate_decay=learning_rate_d  
ecay,  
                      reg=0.0, verbose=True)  
  
    # Predict on the validation set  
    val_acc = (net.predict(X_val) == y_val).mean()  
  
    print('For parameters: ', parameters_matrix[i])  
    print('Validation accuracy: ', val_acc)  
  
    #Store the validation accuracies  
  
    val_acc_vec[i] = val_acc  
  
    #Save the best model and stats  
  
    if val_acc > best_val_acc:  
  
        best_val_acc = val_acc
```

```
best_2layer_sigmoid = net  
best_stats = stats
```

```
#Print the best combination of trial parameters:
```

```
ind = np.argmax(val_acc_vec)
```

```
print('The best combination of parameters is: learning rate = ',parameters_mat  
rix[ind,0], ', hidden layer size = ',int(parameters_matrix[ind,1]), ', batch s  
ize = ',int(parameters_matrix[ind,2]))  
print('For this combination, the validation accuracy is ', val_acc_vec[ind])
```

iteration 0 / 1000: loss 2.394454  
iteration 100 / 1000: loss 1.964292  
iteration 200 / 1000: loss 1.812530  
iteration 300 / 1000: loss 1.871928  
iteration 400 / 1000: loss 1.753414  
iteration 500 / 1000: loss 2.034388  
iteration 600 / 1000: loss 1.926047  
iteration 700 / 1000: loss 1.773440  
iteration 800 / 1000: loss 1.973519  
iteration 900 / 1000: loss 1.611529  
For parameters: [ 0.1 20. 50. ]  
Validation accuracy: 0.401  
iteration 0 / 1000: loss 2.349226  
iteration 100 / 1000: loss 1.873230  
iteration 200 / 1000: loss 1.881829  
iteration 300 / 1000: loss 1.686167  
iteration 400 / 1000: loss 1.657714  
iteration 500 / 1000: loss 1.769798  
iteration 600 / 1000: loss 1.679894  
iteration 700 / 1000: loss 1.534242  
iteration 800 / 1000: loss 1.637482  
iteration 900 / 1000: loss 1.567340  
For parameters: [ 0.1 50. 100. ]  
Validation accuracy: 0.42  
iteration 0 / 1000: loss 2.570874  
iteration 100 / 1000: loss 1.810058  
iteration 200 / 1000: loss 1.679592  
iteration 300 / 1000: loss 1.707853  
iteration 400 / 1000: loss 1.604382  
iteration 500 / 1000: loss 1.711257  
iteration 600 / 1000: loss 1.710577  
iteration 700 / 1000: loss 1.637946  
iteration 800 / 1000: loss 1.621874  
iteration 900 / 1000: loss 1.653794  
For parameters: [1.e-01 8.e+01 2.e+02]  
Validation accuracy: 0.436  
iteration 0 / 1000: loss 2.465977  
iteration 100 / 1000: loss 1.940992  
iteration 200 / 1000: loss 1.906685  
iteration 300 / 1000: loss 1.810134  
iteration 400 / 1000: loss 1.861575  
iteration 500 / 1000: loss 2.028587  
iteration 600 / 1000: loss 1.566310  
iteration 700 / 1000: loss 1.476782  
iteration 800 / 1000: loss 1.578780  
iteration 900 / 1000: loss 1.679164  
For parameters: [ 0.1 100. 50. ]  
Validation accuracy: 0.428  
iteration 0 / 1000: loss 2.328111  
iteration 100 / 1000: loss 1.752327  
iteration 200 / 1000: loss 1.696437  
iteration 300 / 1000: loss 1.724829  
iteration 400 / 1000: loss 1.835532  
iteration 500 / 1000: loss 1.592002  
iteration 600 / 1000: loss 1.707046  
iteration 700 / 1000: loss 1.815531  
iteration 800 / 1000: loss 1.791658

iteration 900 / 1000: loss 1.658150  
For parameters: [1.0e-01 1.2e+02 1.0e+02]  
Validation accuracy: 0.416  
iteration 0 / 1000: loss 2.536106  
iteration 100 / 1000: loss 1.886215  
iteration 200 / 1000: loss 1.863925  
iteration 300 / 1000: loss 1.769263  
iteration 400 / 1000: loss 1.754592  
iteration 500 / 1000: loss 1.670776  
iteration 600 / 1000: loss 1.803482  
iteration 700 / 1000: loss 1.722791  
iteration 800 / 1000: loss 1.567201  
iteration 900 / 1000: loss 1.625860  
For parameters: [1.e-01 2.e+01 2.e+02]  
Validation accuracy: 0.402  
iteration 0 / 1000: loss 2.420396  
iteration 100 / 1000: loss 1.794949  
iteration 200 / 1000: loss 1.744375  
iteration 300 / 1000: loss 1.741356  
iteration 400 / 1000: loss 1.634575  
iteration 500 / 1000: loss 1.591375  
iteration 600 / 1000: loss 1.441955  
iteration 700 / 1000: loss 1.763750  
iteration 800 / 1000: loss 1.647008  
iteration 900 / 1000: loss 1.649860  
For parameters: [ 0.1 50. 50. ]  
Validation accuracy: 0.403  
iteration 0 / 1000: loss 2.323143  
iteration 100 / 1000: loss 1.797190  
iteration 200 / 1000: loss 1.817047  
iteration 300 / 1000: loss 1.858731  
iteration 400 / 1000: loss 1.613066  
iteration 500 / 1000: loss 1.749125  
iteration 600 / 1000: loss 1.773327  
iteration 700 / 1000: loss 1.738443  
iteration 800 / 1000: loss 1.649997  
iteration 900 / 1000: loss 1.555992  
For parameters: [ 0.1 80. 100. ]  
Validation accuracy: 0.423  
iteration 0 / 1000: loss 2.488553  
iteration 100 / 1000: loss 1.757147  
iteration 200 / 1000: loss 1.788353  
iteration 300 / 1000: loss 1.831604  
iteration 400 / 1000: loss 1.739132  
iteration 500 / 1000: loss 1.644011  
iteration 600 / 1000: loss 1.656923  
iteration 700 / 1000: loss 1.537302  
iteration 800 / 1000: loss 1.562717  
iteration 900 / 1000: loss 1.671345  
For parameters: [1.e-01 1.e+02 2.e+02]  
Validation accuracy: 0.415  
iteration 0 / 1000: loss 2.459715  
iteration 100 / 1000: loss 1.703319  
iteration 200 / 1000: loss 2.061390  
iteration 300 / 1000: loss 1.766120  
iteration 400 / 1000: loss 1.917015  
iteration 500 / 1000: loss 1.632332

iteration 600 / 1000: loss 1.933932  
iteration 700 / 1000: loss 1.596196  
iteration 800 / 1000: loss 1.711755  
iteration 900 / 1000: loss 1.656928  
For parameters: [1.0e-01 1.2e+02 5.0e+01]  
Validation accuracy: 0.413  
iteration 0 / 1000: loss 2.487347  
iteration 100 / 1000: loss 1.942118  
iteration 200 / 1000: loss 1.911423  
iteration 300 / 1000: loss 1.845727  
iteration 400 / 1000: loss 1.856512  
iteration 500 / 1000: loss 1.780575  
iteration 600 / 1000: loss 1.678178  
iteration 700 / 1000: loss 1.733185  
iteration 800 / 1000: loss 1.649580  
iteration 900 / 1000: loss 1.754960  
For parameters: [ 0.1 20. 100. ]  
Validation accuracy: 0.424  
iteration 0 / 1000: loss 2.406264  
iteration 100 / 1000: loss 1.820504  
iteration 200 / 1000: loss 1.799362  
iteration 300 / 1000: loss 1.744249  
iteration 400 / 1000: loss 1.750998  
iteration 500 / 1000: loss 1.575253  
iteration 600 / 1000: loss 1.722688  
iteration 700 / 1000: loss 1.661904  
iteration 800 / 1000: loss 1.515969  
iteration 900 / 1000: loss 1.555201  
For parameters: [1.e-01 5.e+01 2.e+02]  
Validation accuracy: 0.422  
iteration 0 / 1000: loss 2.420010  
iteration 100 / 1000: loss 2.016294  
iteration 200 / 1000: loss 1.645791  
iteration 300 / 1000: loss 1.920713  
iteration 400 / 1000: loss 1.582625  
iteration 500 / 1000: loss 1.853434  
iteration 600 / 1000: loss 1.717103  
iteration 700 / 1000: loss 1.844395  
iteration 800 / 1000: loss 1.650734  
iteration 900 / 1000: loss 1.821369  
For parameters: [ 0.1 80. 50. ]  
Validation accuracy: 0.426  
iteration 0 / 1000: loss 2.487527  
iteration 100 / 1000: loss 1.918807  
iteration 200 / 1000: loss 1.560309  
iteration 300 / 1000: loss 1.634781  
iteration 400 / 1000: loss 1.819056  
iteration 500 / 1000: loss 1.778763  
iteration 600 / 1000: loss 1.759877  
iteration 700 / 1000: loss 1.539222  
iteration 800 / 1000: loss 1.625968  
iteration 900 / 1000: loss 1.633689  
For parameters: [ 0.1 100. 100. ]  
Validation accuracy: 0.436  
iteration 0 / 1000: loss 2.385669  
iteration 100 / 1000: loss 1.783753  
iteration 200 / 1000: loss 1.838120

iteration 300 / 1000: loss 1.616081  
iteration 400 / 1000: loss 1.638918  
iteration 500 / 1000: loss 1.682697  
iteration 600 / 1000: loss 1.565685  
iteration 700 / 1000: loss 1.740723  
iteration 800 / 1000: loss 1.699835  
iteration 900 / 1000: loss 1.649610  
For parameters: [1.0e-01 1.2e+02 2.0e+02]  
Validation accuracy: 0.424  
iteration 0 / 1000: loss 2.455716  
iteration 100 / 1000: loss 2.065960  
iteration 200 / 1000: loss 1.855749  
iteration 300 / 1000: loss 2.001173  
iteration 400 / 1000: loss 1.777433  
iteration 500 / 1000: loss 1.789087  
iteration 600 / 1000: loss 1.730583  
iteration 700 / 1000: loss 1.703707  
iteration 800 / 1000: loss 1.752186  
iteration 900 / 1000: loss 1.734850  
For parameters: [ 0.05 20. 50. ]  
Validation accuracy: 0.408  
iteration 0 / 1000: loss 2.420198  
iteration 100 / 1000: loss 1.806616  
iteration 200 / 1000: loss 1.746601  
iteration 300 / 1000: loss 1.839837  
iteration 400 / 1000: loss 1.718981  
iteration 500 / 1000: loss 1.633449  
iteration 600 / 1000: loss 1.711333  
iteration 700 / 1000: loss 1.647152  
iteration 800 / 1000: loss 1.836232  
iteration 900 / 1000: loss 1.853690  
For parameters: [5.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.418  
iteration 0 / 1000: loss 2.373002  
iteration 100 / 1000: loss 1.904264  
iteration 200 / 1000: loss 1.760397  
iteration 300 / 1000: loss 1.734230  
iteration 400 / 1000: loss 1.643912  
iteration 500 / 1000: loss 1.779889  
iteration 600 / 1000: loss 1.701976  
iteration 700 / 1000: loss 1.780676  
iteration 800 / 1000: loss 1.657758  
iteration 900 / 1000: loss 1.732420  
For parameters: [5.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.412  
iteration 0 / 1000: loss 2.282326  
iteration 100 / 1000: loss 1.935968  
iteration 200 / 1000: loss 1.755951  
iteration 300 / 1000: loss 1.693249  
iteration 400 / 1000: loss 1.778067  
iteration 500 / 1000: loss 1.790226  
iteration 600 / 1000: loss 1.809796  
iteration 700 / 1000: loss 1.933012  
iteration 800 / 1000: loss 1.809176  
iteration 900 / 1000: loss 1.881093  
For parameters: [5.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.423

iteration 0 / 1000: loss 2.400095  
iteration 100 / 1000: loss 1.831414  
iteration 200 / 1000: loss 1.800341  
iteration 300 / 1000: loss 1.647824  
iteration 400 / 1000: loss 1.650170  
iteration 500 / 1000: loss 1.744460  
iteration 600 / 1000: loss 1.689547  
iteration 700 / 1000: loss 1.907491  
iteration 800 / 1000: loss 1.786880  
iteration 900 / 1000: loss 1.669094  
For parameters: [5.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.418  
iteration 0 / 1000: loss 2.476107  
iteration 100 / 1000: loss 1.960507  
iteration 200 / 1000: loss 1.858606  
iteration 300 / 1000: loss 1.862437  
iteration 400 / 1000: loss 1.893980  
iteration 500 / 1000: loss 1.803977  
iteration 600 / 1000: loss 1.747360  
iteration 700 / 1000: loss 1.792101  
iteration 800 / 1000: loss 1.754027  
iteration 900 / 1000: loss 1.756797  
For parameters: [5.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.394  
iteration 0 / 1000: loss 2.399920  
iteration 100 / 1000: loss 1.914639  
iteration 200 / 1000: loss 1.911489  
iteration 300 / 1000: loss 1.784970  
iteration 400 / 1000: loss 1.713829  
iteration 500 / 1000: loss 1.696032  
iteration 600 / 1000: loss 1.729354  
iteration 700 / 1000: loss 1.627333  
iteration 800 / 1000: loss 1.813409  
iteration 900 / 1000: loss 1.592761  
For parameters: [ 0.05 50. 50. ]  
Validation accuracy: 0.428  
iteration 0 / 1000: loss 2.353592  
iteration 100 / 1000: loss 2.019339  
iteration 200 / 1000: loss 1.909803  
iteration 300 / 1000: loss 1.701646  
iteration 400 / 1000: loss 1.786880  
iteration 500 / 1000: loss 1.758090  
iteration 600 / 1000: loss 1.703448  
iteration 700 / 1000: loss 1.595363  
iteration 800 / 1000: loss 1.641676  
iteration 900 / 1000: loss 1.738780  
For parameters: [5.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.42  
iteration 0 / 1000: loss 2.345795  
iteration 100 / 1000: loss 1.815835  
iteration 200 / 1000: loss 1.846279  
iteration 300 / 1000: loss 1.794705  
iteration 400 / 1000: loss 1.733719  
iteration 500 / 1000: loss 1.798558  
iteration 600 / 1000: loss 1.620289  
iteration 700 / 1000: loss 1.760604  
iteration 800 / 1000: loss 1.761495



iteration 900 / 1000: loss 1.677075  
For parameters: [5.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.414  
iteration 0 / 1000: loss 2.469060  
iteration 100 / 1000: loss 1.930331  
iteration 200 / 1000: loss 1.770706  
iteration 300 / 1000: loss 1.733105  
iteration 400 / 1000: loss 1.754774  
iteration 500 / 1000: loss 1.604754  
iteration 600 / 1000: loss 1.831471  
iteration 700 / 1000: loss 1.644048  
iteration 800 / 1000: loss 1.546032  
iteration 900 / 1000: loss 1.649707  
For parameters: [5.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.402  
iteration 0 / 1000: loss 2.479579  
iteration 100 / 1000: loss 1.944570  
iteration 200 / 1000: loss 1.934310  
iteration 300 / 1000: loss 1.799130  
iteration 400 / 1000: loss 1.735172  
iteration 500 / 1000: loss 1.762588  
iteration 600 / 1000: loss 1.790943  
iteration 700 / 1000: loss 1.840967  
iteration 800 / 1000: loss 1.715145  
iteration 900 / 1000: loss 1.762889  
For parameters: [5.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.41  
iteration 0 / 1000: loss 2.416241  
iteration 100 / 1000: loss 1.945121  
iteration 200 / 1000: loss 1.776174  
iteration 300 / 1000: loss 1.901165  
iteration 400 / 1000: loss 1.715005  
iteration 500 / 1000: loss 1.770175  
iteration 600 / 1000: loss 1.740622  
iteration 700 / 1000: loss 1.701591  
iteration 800 / 1000: loss 1.671792  
iteration 900 / 1000: loss 1.693164  
For parameters: [5.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.42  
iteration 0 / 1000: loss 2.375107  
iteration 100 / 1000: loss 2.043450  
iteration 200 / 1000: loss 1.928399  
iteration 300 / 1000: loss 1.785400  
iteration 400 / 1000: loss 1.687530  
iteration 500 / 1000: loss 1.859464  
iteration 600 / 1000: loss 1.617911  
iteration 700 / 1000: loss 1.607465  
iteration 800 / 1000: loss 1.714931  
iteration 900 / 1000: loss 1.572047  
For parameters: [5.e-02 8.e+01 5.e+01]  
Validation accuracy: 0.416  
iteration 0 / 1000: loss 2.374696  
iteration 100 / 1000: loss 1.871647  
iteration 200 / 1000: loss 1.857009  
iteration 300 / 1000: loss 1.818499  
iteration 400 / 1000: loss 1.694697  
iteration 500 / 1000: loss 1.813534

iteration 600 / 1000: loss 1.755415  
iteration 700 / 1000: loss 1.688975  
iteration 800 / 1000: loss 1.814380  
iteration 900 / 1000: loss 1.769124  
For parameters: [5.e-02 1.e+02 1.e+02]  
Validation accuracy: 0.406  
iteration 0 / 1000: loss 2.412279  
iteration 100 / 1000: loss 1.871401  
iteration 200 / 1000: loss 1.867153  
iteration 300 / 1000: loss 1.773904  
iteration 400 / 1000: loss 1.752046  
iteration 500 / 1000: loss 1.672935  
iteration 600 / 1000: loss 1.754616  
iteration 700 / 1000: loss 1.687278  
iteration 800 / 1000: loss 1.712759  
iteration 900 / 1000: loss 1.603706  
For parameters: [5.0e-02 1.2e+02 2.0e+02]  
Validation accuracy: 0.424  
iteration 0 / 1000: loss 2.350478  
iteration 100 / 1000: loss 2.026968  
iteration 200 / 1000: loss 1.971370  
iteration 300 / 1000: loss 1.808695  
iteration 400 / 1000: loss 1.902631  
iteration 500 / 1000: loss 1.901005  
iteration 600 / 1000: loss 1.766567  
iteration 700 / 1000: loss 1.691466  
iteration 800 / 1000: loss 1.998961  
iteration 900 / 1000: loss 1.698829  
For parameters: [3.e-02 2.e+01 5.e+01]  
Validation accuracy: 0.386  
iteration 0 / 1000: loss 2.340535  
iteration 100 / 1000: loss 1.998176  
iteration 200 / 1000: loss 1.885393  
iteration 300 / 1000: loss 1.939431  
iteration 400 / 1000: loss 1.816331  
iteration 500 / 1000: loss 1.869912  
iteration 600 / 1000: loss 1.743681  
iteration 700 / 1000: loss 1.927129  
iteration 800 / 1000: loss 1.694740  
iteration 900 / 1000: loss 1.799152  
For parameters: [3.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.399  
iteration 0 / 1000: loss 2.482252  
iteration 100 / 1000: loss 1.994098  
iteration 200 / 1000: loss 1.873706  
iteration 300 / 1000: loss 1.834685  
iteration 400 / 1000: loss 1.867152  
iteration 500 / 1000: loss 1.679867  
iteration 600 / 1000: loss 1.844308  
iteration 700 / 1000: loss 1.619347  
iteration 800 / 1000: loss 1.776429  
iteration 900 / 1000: loss 1.610273  
For parameters: [3.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.416  
iteration 0 / 1000: loss 2.486651  
iteration 100 / 1000: loss 1.908359  
iteration 200 / 1000: loss 1.848283

iteration 300 / 1000: loss 1.599898  
iteration 400 / 1000: loss 1.836110  
iteration 500 / 1000: loss 1.751169  
iteration 600 / 1000: loss 1.821119  
iteration 700 / 1000: loss 1.906609  
iteration 800 / 1000: loss 1.587478  
iteration 900 / 1000: loss 1.576083  
For parameters: [3.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.409  
iteration 0 / 1000: loss 2.347113  
iteration 100 / 1000: loss 1.822641  
iteration 200 / 1000: loss 1.899611  
iteration 300 / 1000: loss 1.982349  
iteration 400 / 1000: loss 1.871965  
iteration 500 / 1000: loss 1.705308  
iteration 600 / 1000: loss 1.883409  
iteration 700 / 1000: loss 1.724040  
iteration 800 / 1000: loss 1.695064  
iteration 900 / 1000: loss 1.711206  
For parameters: [3.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.406  
iteration 0 / 1000: loss 2.395162  
iteration 100 / 1000: loss 2.006948  
iteration 200 / 1000: loss 1.956431  
iteration 300 / 1000: loss 1.987345  
iteration 400 / 1000: loss 1.903888  
iteration 500 / 1000: loss 1.850027  
iteration 600 / 1000: loss 1.930595  
iteration 700 / 1000: loss 1.767411  
iteration 800 / 1000: loss 1.826727  
iteration 900 / 1000: loss 1.897840  
For parameters: [3.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.398  
iteration 0 / 1000: loss 2.312706  
iteration 100 / 1000: loss 1.915991  
iteration 200 / 1000: loss 1.988096  
iteration 300 / 1000: loss 1.854926  
iteration 400 / 1000: loss 1.940358  
iteration 500 / 1000: loss 1.912949  
iteration 600 / 1000: loss 1.756509  
iteration 700 / 1000: loss 1.827021  
iteration 800 / 1000: loss 1.795509  
iteration 900 / 1000: loss 1.707773  
For parameters: [3.e-02 5.e+01 5.e+01]  
Validation accuracy: 0.407  
iteration 0 / 1000: loss 2.464060  
iteration 100 / 1000: loss 2.008343  
iteration 200 / 1000: loss 1.703615  
iteration 300 / 1000: loss 1.906412  
iteration 400 / 1000: loss 1.775457  
iteration 500 / 1000: loss 1.727726  
iteration 600 / 1000: loss 1.742025  
iteration 700 / 1000: loss 1.818305  
iteration 800 / 1000: loss 1.837705  
iteration 900 / 1000: loss 1.795398  
For parameters: [3.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.415

iteration 0 / 1000: loss 2.453507  
iteration 100 / 1000: loss 1.888267  
iteration 200 / 1000: loss 1.846556  
iteration 300 / 1000: loss 1.869506  
iteration 400 / 1000: loss 1.778286  
iteration 500 / 1000: loss 1.812916  
iteration 600 / 1000: loss 1.840019  
iteration 700 / 1000: loss 1.590447  
iteration 800 / 1000: loss 1.676140  
iteration 900 / 1000: loss 1.559915  
For parameters: [3.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.402  
iteration 0 / 1000: loss 2.496868  
iteration 100 / 1000: loss 1.960003  
iteration 200 / 1000: loss 1.817180  
iteration 300 / 1000: loss 2.041408  
iteration 400 / 1000: loss 1.822522  
iteration 500 / 1000: loss 1.620768  
iteration 600 / 1000: loss 1.958609  
iteration 700 / 1000: loss 1.641498  
iteration 800 / 1000: loss 1.776475  
iteration 900 / 1000: loss 1.828980  
For parameters: [3.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.393  
iteration 0 / 1000: loss 2.409908  
iteration 100 / 1000: loss 2.035423  
iteration 200 / 1000: loss 1.920146  
iteration 300 / 1000: loss 2.013233  
iteration 400 / 1000: loss 1.809811  
iteration 500 / 1000: loss 1.958199  
iteration 600 / 1000: loss 1.734700  
iteration 700 / 1000: loss 1.779192  
iteration 800 / 1000: loss 1.697347  
iteration 900 / 1000: loss 1.835212  
For parameters: [3.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.394  
iteration 0 / 1000: loss 2.430125  
iteration 100 / 1000: loss 1.980472  
iteration 200 / 1000: loss 1.894741  
iteration 300 / 1000: loss 1.876238  
iteration 400 / 1000: loss 1.790473  
iteration 500 / 1000: loss 1.768539  
iteration 600 / 1000: loss 1.786717  
iteration 700 / 1000: loss 1.751591  
iteration 800 / 1000: loss 1.620372  
iteration 900 / 1000: loss 1.717104  
For parameters: [3.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.408  
iteration 0 / 1000: loss 2.329865  
iteration 100 / 1000: loss 1.922829  
iteration 200 / 1000: loss 1.866684  
iteration 300 / 1000: loss 1.862252  
iteration 400 / 1000: loss 1.809355  
iteration 500 / 1000: loss 1.688395  
iteration 600 / 1000: loss 1.758873  
iteration 700 / 1000: loss 1.601702  
iteration 800 / 1000: loss 1.975365

```
iteration 900 / 1000: loss 1.793599
For parameters: [3.e-02 8.e+01 5.e+01]
Validation accuracy: 0.411
iteration 0 / 1000: loss 2.532448
iteration 100 / 1000: loss 1.974171
iteration 200 / 1000: loss 1.714697
iteration 300 / 1000: loss 1.844785
iteration 400 / 1000: loss 1.886980
iteration 500 / 1000: loss 1.673030
iteration 600 / 1000: loss 1.871183
iteration 700 / 1000: loss 1.729487
iteration 800 / 1000: loss 1.915398
iteration 900 / 1000: loss 1.631230
For parameters: [3.e-02 1.e+02 1.e+02]
Validation accuracy: 0.415
iteration 0 / 1000: loss 2.345802
iteration 100 / 1000: loss 1.854969
iteration 200 / 1000: loss 1.849685
iteration 300 / 1000: loss 1.862382
iteration 400 / 1000: loss 1.733506
iteration 500 / 1000: loss 1.847152
iteration 600 / 1000: loss 1.769355
iteration 700 / 1000: loss 1.759110
iteration 800 / 1000: loss 1.783592
iteration 900 / 1000: loss 1.725988
For parameters: [3.0e-02 1.2e+02 2.0e+02]
Validation accuracy: 0.406
The best combination of parameters is: learning rate = 0.1 , hidden layer si
ze = 80 , batch size = 200
For this combination, the validation accuracy is 0.436
```

## Three-layer Relu Activation Network

In [16]: best\_3layer\_relu = **None** # store the best model into this

```
#####  
###  
# TODO: Tune hyperparameters using the validation set. Store your best trained  
#  
# model in best_3layer_relu.  
#  
#####  
###  
  
#Initialize "best" validation accuracy  
  
best_val_acc = 0  
  
#Initialize vector to store all validation accuracies  
  
val_acc_vec = np.zeros(len(parameters_matrix))  
  
#Cycle through the matrix of parameters  
  
for i in range(len(parameters_matrix)):  
  
    input_size = 32 * 32 * 3  
    num_layers = 3  
    hidden_size = int(parameters_matrix[i,1])  
    hidden_sizes = [hidden_size]*(num_layers-1)  
    num_classes = 10  
    learning_rate = parameters_matrix[i,0]  
    learning_rate_decay = 0.95  
    batch_size = int(parameters_matrix[i,2])  
    net = NeuralNetwork(input_size, hidden_sizes, num_classes, num_layers, nonli  
nearity='relu')  
  
    # Train the network  
    stats = net.train(X_train, y_train, X_val, y_val,  
                      num_iters=1000, batch_size=batch_size,  
                      learning_rate=learning_rate, learning_rate_decay=learning_rate_d  
ecay,  
                      reg=0.0, verbose=True)  
  
    # Predict on the validation set  
    val_acc = (net.predict(X_val) == y_val).mean()  
  
    print('For parameters: ', parameters_matrix[i])  
    print('Validation accuracy: ', val_acc)  
  
    #Store the validation accuracies  
  
    val_acc_vec[i] = val_acc  
  
    #Save the best model and stats  
  
    if val_acc > best_val_acc:  
  
        best_val_acc = val_acc
```

```
best_3layer_relu = net  
best_stats = stats
```

```
#Print the best combination of trial parameters:
```

```
ind = np.argmax(val_acc_vec)
```

```
print('The best combination of parameters is: learning rate = ',parameters_mat  
rix[ind,0], ', hidden layer size = ',int(parameters_matrix[ind,1]), ', batch s  
ize = ',int(parameters_matrix[ind,2]))  
print('For this combination, the validation accuracy is ', val_acc_vec[ind])
```

iteration 0 / 1000: loss 2.488548  
iteration 100 / 1000: loss 2.149827  
iteration 200 / 1000: loss 1.861412  
iteration 300 / 1000: loss 1.573523  
iteration 400 / 1000: loss 1.660306  
iteration 500 / 1000: loss 1.895783  
iteration 600 / 1000: loss 1.484960  
iteration 700 / 1000: loss 1.507817  
iteration 800 / 1000: loss 1.812805  
iteration 900 / 1000: loss 1.652591  
For parameters: [ 0.1 20. 50. ]  
Validation accuracy: 0.398  
iteration 0 / 1000: loss 2.506757  
iteration 100 / 1000: loss 1.790607  
iteration 200 / 1000: loss 1.535027  
iteration 300 / 1000: loss 1.468571  
iteration 400 / 1000: loss 1.512914  
iteration 500 / 1000: loss 1.590604  
iteration 600 / 1000: loss 1.384771  
iteration 700 / 1000: loss 1.651912  
iteration 800 / 1000: loss 1.385913  
iteration 900 / 1000: loss 1.587206  
For parameters: [ 0.1 50. 100. ]  
Validation accuracy: 0.459  
iteration 0 / 1000: loss 2.356532  
iteration 100 / 1000: loss 1.686930  
iteration 200 / 1000: loss 1.563930  
iteration 300 / 1000: loss 1.420197  
iteration 400 / 1000: loss 1.643392  
iteration 500 / 1000: loss 1.403180  
iteration 600 / 1000: loss 1.398474  
iteration 700 / 1000: loss 1.289954  
iteration 800 / 1000: loss 1.273592  
iteration 900 / 1000: loss 1.416335  
For parameters: [1.e-01 8.e+01 2.e+02]  
Validation accuracy: 0.49  
iteration 0 / 1000: loss 2.427977  
iteration 100 / 1000: loss 1.795480  
iteration 200 / 1000: loss 1.683010  
iteration 300 / 1000: loss 1.726709  
iteration 400 / 1000: loss 1.568488  
iteration 500 / 1000: loss 1.400907  
iteration 600 / 1000: loss 1.580139  
iteration 700 / 1000: loss 1.256653  
iteration 800 / 1000: loss 1.615023  
iteration 900 / 1000: loss 1.370663  
For parameters: [ 0.1 100. 50. ]  
Validation accuracy: 0.447  
iteration 0 / 1000: loss 2.368715  
iteration 100 / 1000: loss 1.690701  
iteration 200 / 1000: loss 1.571183  
iteration 300 / 1000: loss 1.525240  
iteration 400 / 1000: loss 1.494508  
iteration 500 / 1000: loss 1.591508  
iteration 600 / 1000: loss 1.668100  
iteration 700 / 1000: loss 1.352411  
iteration 800 / 1000: loss 1.368516



iteration 900 / 1000: loss 1.499359  
For parameters: [1.0e-01 1.2e+02 1.0e+02]  
Validation accuracy: 0.503  
iteration 0 / 1000: loss 2.366797  
iteration 100 / 1000: loss 1.871385  
iteration 200 / 1000: loss 1.643839  
iteration 300 / 1000: loss 1.508153  
iteration 400 / 1000: loss 1.581833  
iteration 500 / 1000: loss 1.568563  
iteration 600 / 1000: loss 1.560668  
iteration 700 / 1000: loss 1.669056  
iteration 800 / 1000: loss 1.464973  
iteration 900 / 1000: loss 1.552902  
For parameters: [1.e-01 2.e+01 2.e+02]  
Validation accuracy: 0.441  
iteration 0 / 1000: loss 2.410853  
iteration 100 / 1000: loss 1.771586  
iteration 200 / 1000: loss 2.127667  
iteration 300 / 1000: loss 1.591124  
iteration 400 / 1000: loss 1.640894  
iteration 500 / 1000: loss 1.707135  
iteration 600 / 1000: loss 1.651658  
iteration 700 / 1000: loss 1.807850  
iteration 800 / 1000: loss 1.491906  
iteration 900 / 1000: loss 1.324043  
For parameters: [ 0.1 50. 50. ]  
Validation accuracy: 0.437  
iteration 0 / 1000: loss 2.396920  
iteration 100 / 1000: loss 1.857339  
iteration 200 / 1000: loss 1.536795  
iteration 300 / 1000: loss 1.698601  
iteration 400 / 1000: loss 1.466450  
iteration 500 / 1000: loss 1.408033  
iteration 600 / 1000: loss 1.547902  
iteration 700 / 1000: loss 1.691444  
iteration 800 / 1000: loss 1.418001  
iteration 900 / 1000: loss 1.277043  
For parameters: [ 0.1 80. 100. ]  
Validation accuracy: 0.465  
iteration 0 / 1000: loss 2.403956  
iteration 100 / 1000: loss 1.732292  
iteration 200 / 1000: loss 1.667944  
iteration 300 / 1000: loss 1.454914  
iteration 400 / 1000: loss 1.429880  
iteration 500 / 1000: loss 1.418885  
iteration 600 / 1000: loss 1.542091  
iteration 700 / 1000: loss 1.366346  
iteration 800 / 1000: loss 1.290044  
iteration 900 / 1000: loss 1.229839  
For parameters: [1.e-01 1.e+02 2.e+02]  
Validation accuracy: 0.507  
iteration 0 / 1000: loss 2.387579  
iteration 100 / 1000: loss 1.828161  
iteration 200 / 1000: loss 1.648334  
iteration 300 / 1000: loss 1.788637  
iteration 400 / 1000: loss 1.592732  
iteration 500 / 1000: loss 1.225307

iteration 600 / 1000: loss 1.482539  
iteration 700 / 1000: loss 1.486220  
iteration 800 / 1000: loss 1.460764  
iteration 900 / 1000: loss 1.558186  
For parameters: [1.0e-01 1.2e+02 5.0e+01]  
Validation accuracy: 0.45  
iteration 0 / 1000: loss 2.341265  
iteration 100 / 1000: loss 1.834678  
iteration 200 / 1000: loss 1.897337  
iteration 300 / 1000: loss 1.619981  
iteration 400 / 1000: loss 1.547761  
iteration 500 / 1000: loss 1.729191  
iteration 600 / 1000: loss 1.734293  
iteration 700 / 1000: loss 1.533691  
iteration 800 / 1000: loss 1.534677  
iteration 900 / 1000: loss 1.280966  
For parameters: [ 0.1 20. 100. ]  
Validation accuracy: 0.441  
iteration 0 / 1000: loss 2.379973  
iteration 100 / 1000: loss 1.711375  
iteration 200 / 1000: loss 1.580873  
iteration 300 / 1000: loss 1.512723  
iteration 400 / 1000: loss 1.486065  
iteration 500 / 1000: loss 1.475074  
iteration 600 / 1000: loss 1.567917  
iteration 700 / 1000: loss 1.465371  
iteration 800 / 1000: loss 1.338320  
iteration 900 / 1000: loss 1.419387  
For parameters: [1.e-01 5.e+01 2.e+02]  
Validation accuracy: 0.492  
iteration 0 / 1000: loss 2.474014  
iteration 100 / 1000: loss 1.787571  
iteration 200 / 1000: loss 1.740740  
iteration 300 / 1000: loss 1.780470  
iteration 400 / 1000: loss 1.746044  
iteration 500 / 1000: loss 1.864854  
iteration 600 / 1000: loss 1.773115  
iteration 700 / 1000: loss 1.789275  
iteration 800 / 1000: loss 1.453697  
iteration 900 / 1000: loss 1.689362  
For parameters: [ 0.1 80. 50. ]  
Validation accuracy: 0.44  
iteration 0 / 1000: loss 2.496058  
iteration 100 / 1000: loss 1.863945  
iteration 200 / 1000: loss 1.682928  
iteration 300 / 1000: loss 1.536935  
iteration 400 / 1000: loss 1.511178  
iteration 500 / 1000: loss 1.430720  
iteration 600 / 1000: loss 1.507224  
iteration 700 / 1000: loss 1.529159  
iteration 800 / 1000: loss 1.497969  
iteration 900 / 1000: loss 1.380143  
For parameters: [ 0.1 100. 100. ]  
Validation accuracy: 0.475  
iteration 0 / 1000: loss 2.435805  
iteration 100 / 1000: loss 1.635913  
iteration 200 / 1000: loss 1.724216

iteration 300 / 1000: loss 1.564232  
iteration 400 / 1000: loss 1.522211  
iteration 500 / 1000: loss 1.524997  
iteration 600 / 1000: loss 1.324900  
iteration 700 / 1000: loss 1.413647  
iteration 800 / 1000: loss 1.249312  
iteration 900 / 1000: loss 1.269321  
For parameters: [1.0e-01 1.2e+02 2.0e+02]  
Validation accuracy: 0.477  
iteration 0 / 1000: loss 2.483989  
iteration 100 / 1000: loss 1.850517  
iteration 200 / 1000: loss 1.908296  
iteration 300 / 1000: loss 1.808306  
iteration 400 / 1000: loss 1.622488  
iteration 500 / 1000: loss 1.725570  
iteration 600 / 1000: loss 1.549689  
iteration 700 / 1000: loss 1.715273  
iteration 800 / 1000: loss 1.739351  
iteration 900 / 1000: loss 1.736076  
For parameters: [ 0.05 20. 50. ]  
Validation accuracy: 0.422  
iteration 0 / 1000: loss 2.478430  
iteration 100 / 1000: loss 1.807153  
iteration 200 / 1000: loss 1.873974  
iteration 300 / 1000: loss 1.626945  
iteration 400 / 1000: loss 1.464768  
iteration 500 / 1000: loss 1.438573  
iteration 600 / 1000: loss 1.444805  
iteration 700 / 1000: loss 1.695742  
iteration 800 / 1000: loss 1.380290  
iteration 900 / 1000: loss 1.434087  
For parameters: [5.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.466  
iteration 0 / 1000: loss 2.467992  
iteration 100 / 1000: loss 1.655150  
iteration 200 / 1000: loss 1.604073  
iteration 300 / 1000: loss 1.459405  
iteration 400 / 1000: loss 1.435752  
iteration 500 / 1000: loss 1.382560  
iteration 600 / 1000: loss 1.352078  
iteration 700 / 1000: loss 1.266220  
iteration 800 / 1000: loss 1.358554  
iteration 900 / 1000: loss 1.434934  
For parameters: [5.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.481  
iteration 0 / 1000: loss 2.389696  
iteration 100 / 1000: loss 2.028125  
iteration 200 / 1000: loss 1.573349  
iteration 300 / 1000: loss 1.424485  
iteration 400 / 1000: loss 1.756657  
iteration 500 / 1000: loss 1.604768  
iteration 600 / 1000: loss 1.816790  
iteration 700 / 1000: loss 1.598568  
iteration 800 / 1000: loss 1.417245  
iteration 900 / 1000: loss 1.768143  
For parameters: [5.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.458

iteration 0 / 1000: loss 2.477916  
iteration 100 / 1000: loss 1.628520  
iteration 200 / 1000: loss 1.738098  
iteration 300 / 1000: loss 1.715109  
iteration 400 / 1000: loss 1.606572  
iteration 500 / 1000: loss 1.390299  
iteration 600 / 1000: loss 1.504065  
iteration 700 / 1000: loss 1.485633  
iteration 800 / 1000: loss 1.484221  
iteration 900 / 1000: loss 1.408944  
For parameters: [5.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.46  
iteration 0 / 1000: loss 2.378349  
iteration 100 / 1000: loss 1.798541  
iteration 200 / 1000: loss 1.741227  
iteration 300 / 1000: loss 1.744580  
iteration 400 / 1000: loss 1.682288  
iteration 500 / 1000: loss 1.612828  
iteration 600 / 1000: loss 1.647541  
iteration 700 / 1000: loss 1.564386  
iteration 800 / 1000: loss 1.521308  
iteration 900 / 1000: loss 1.496103  
For parameters: [5.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.435  
iteration 0 / 1000: loss 2.403335  
iteration 100 / 1000: loss 1.760945  
iteration 200 / 1000: loss 1.788317  
iteration 300 / 1000: loss 1.673200  
iteration 400 / 1000: loss 1.759980  
iteration 500 / 1000: loss 1.639813  
iteration 600 / 1000: loss 1.630916  
iteration 700 / 1000: loss 1.753673  
iteration 800 / 1000: loss 1.867768  
iteration 900 / 1000: loss 1.692217  
For parameters: [ 0.05 50. 50. ]  
Validation accuracy: 0.447  
iteration 0 / 1000: loss 2.461821  
iteration 100 / 1000: loss 1.652415  
iteration 200 / 1000: loss 1.618583  
iteration 300 / 1000: loss 1.562309  
iteration 400 / 1000: loss 1.530729  
iteration 500 / 1000: loss 1.476320  
iteration 600 / 1000: loss 1.678424  
iteration 700 / 1000: loss 1.483633  
iteration 800 / 1000: loss 1.443606  
iteration 900 / 1000: loss 1.396224  
For parameters: [5.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.477  
iteration 0 / 1000: loss 2.352702  
iteration 100 / 1000: loss 1.752401  
iteration 200 / 1000: loss 1.647341  
iteration 300 / 1000: loss 1.613936  
iteration 400 / 1000: loss 1.540483  
iteration 500 / 1000: loss 1.379023  
iteration 600 / 1000: loss 1.421652  
iteration 700 / 1000: loss 1.408681  
iteration 800 / 1000: loss 1.390706

iteration 900 / 1000: loss 1.262250  
For parameters: [5.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.492  
iteration 0 / 1000: loss 2.531880  
iteration 100 / 1000: loss 1.759208  
iteration 200 / 1000: loss 1.777803  
iteration 300 / 1000: loss 1.403432  
iteration 400 / 1000: loss 1.411589  
iteration 500 / 1000: loss 1.783714  
iteration 600 / 1000: loss 1.325228  
iteration 700 / 1000: loss 1.551734  
iteration 800 / 1000: loss 1.339371  
iteration 900 / 1000: loss 1.531675  
For parameters: [5.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.461  
iteration 0 / 1000: loss 2.417943  
iteration 100 / 1000: loss 1.767447  
iteration 200 / 1000: loss 1.631187  
iteration 300 / 1000: loss 1.584551  
iteration 400 / 1000: loss 1.649359  
iteration 500 / 1000: loss 1.602876  
iteration 600 / 1000: loss 1.657365  
iteration 700 / 1000: loss 1.596195  
iteration 800 / 1000: loss 1.555987  
iteration 900 / 1000: loss 1.501673  
For parameters: [5.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.411  
iteration 0 / 1000: loss 2.462928  
iteration 100 / 1000: loss 1.800741  
iteration 200 / 1000: loss 1.657112  
iteration 300 / 1000: loss 1.605263  
iteration 400 / 1000: loss 1.610468  
iteration 500 / 1000: loss 1.592227  
iteration 600 / 1000: loss 1.465059  
iteration 700 / 1000: loss 1.428843  
iteration 800 / 1000: loss 1.554932  
iteration 900 / 1000: loss 1.430861  
For parameters: [5.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.463  
iteration 0 / 1000: loss 2.319795  
iteration 100 / 1000: loss 1.750372  
iteration 200 / 1000: loss 1.892024  
iteration 300 / 1000: loss 1.643956  
iteration 400 / 1000: loss 1.566700  
iteration 500 / 1000: loss 1.775872  
iteration 600 / 1000: loss 1.392110  
iteration 700 / 1000: loss 1.513182  
iteration 800 / 1000: loss 1.620749  
iteration 900 / 1000: loss 1.560164  
For parameters: [5.e-02 8.e+01 5.e+01]  
Validation accuracy: 0.451  
iteration 0 / 1000: loss 2.454141  
iteration 100 / 1000: loss 1.579955  
iteration 200 / 1000: loss 1.642354  
iteration 300 / 1000: loss 1.650825  
iteration 400 / 1000: loss 1.715080  
iteration 500 / 1000: loss 1.457042

iteration 600 / 1000: loss 1.522399  
iteration 700 / 1000: loss 1.478631  
iteration 800 / 1000: loss 1.516866  
iteration 900 / 1000: loss 1.260067  
For parameters: [5.e-02 1.e+02 1.e+02]  
Validation accuracy: 0.495  
iteration 0 / 1000: loss 2.386968  
iteration 100 / 1000: loss 1.596986  
iteration 200 / 1000: loss 1.554853  
iteration 300 / 1000: loss 1.583574  
iteration 400 / 1000: loss 1.402773  
iteration 500 / 1000: loss 1.378430  
iteration 600 / 1000: loss 1.375337  
iteration 700 / 1000: loss 1.315856  
iteration 800 / 1000: loss 1.275970  
iteration 900 / 1000: loss 1.365006  
For parameters: [5.0e-02 1.2e+02 2.0e+02]  
Validation accuracy: 0.501  
iteration 0 / 1000: loss 2.448250  
iteration 100 / 1000: loss 2.110432  
iteration 200 / 1000: loss 1.665533  
iteration 300 / 1000: loss 1.843127  
iteration 400 / 1000: loss 1.807560  
iteration 500 / 1000: loss 1.773654  
iteration 600 / 1000: loss 1.761658  
iteration 700 / 1000: loss 1.285798  
iteration 800 / 1000: loss 1.898557  
iteration 900 / 1000: loss 1.426165  
For parameters: [3.e-02 2.e+01 5.e+01]  
Validation accuracy: 0.408  
iteration 0 / 1000: loss 2.422844  
iteration 100 / 1000: loss 1.886660  
iteration 200 / 1000: loss 1.700300  
iteration 300 / 1000: loss 1.896407  
iteration 400 / 1000: loss 1.371108  
iteration 500 / 1000: loss 1.518603  
iteration 600 / 1000: loss 1.491910  
iteration 700 / 1000: loss 1.587119  
iteration 800 / 1000: loss 1.451214  
iteration 900 / 1000: loss 1.448096  
For parameters: [3.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.461  
iteration 0 / 1000: loss 2.432160  
iteration 100 / 1000: loss 1.759717  
iteration 200 / 1000: loss 1.637465  
iteration 300 / 1000: loss 1.685979  
iteration 400 / 1000: loss 1.483603  
iteration 500 / 1000: loss 1.545669  
iteration 600 / 1000: loss 1.441953  
iteration 700 / 1000: loss 1.512083  
iteration 800 / 1000: loss 1.470915  
iteration 900 / 1000: loss 1.430387  
For parameters: [3.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.488  
iteration 0 / 1000: loss 2.475252  
iteration 100 / 1000: loss 1.811449  
iteration 200 / 1000: loss 1.636489

iteration 300 / 1000: loss 1.640049  
iteration 400 / 1000: loss 1.721024  
iteration 500 / 1000: loss 1.778773  
iteration 600 / 1000: loss 1.373376  
iteration 700 / 1000: loss 1.557382  
iteration 800 / 1000: loss 1.594996  
iteration 900 / 1000: loss 1.253603  
For parameters: [3.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.452  
iteration 0 / 1000: loss 2.435103  
iteration 100 / 1000: loss 1.801877  
iteration 200 / 1000: loss 1.673828  
iteration 300 / 1000: loss 1.692316  
iteration 400 / 1000: loss 1.661874  
iteration 500 / 1000: loss 1.472946  
iteration 600 / 1000: loss 1.600589  
iteration 700 / 1000: loss 1.433157  
iteration 800 / 1000: loss 1.437748  
iteration 900 / 1000: loss 1.228969  
For parameters: [3.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.479  
iteration 0 / 1000: loss 2.347293  
iteration 100 / 1000: loss 1.899067  
iteration 200 / 1000: loss 1.753977  
iteration 300 / 1000: loss 1.684895  
iteration 400 / 1000: loss 1.737257  
iteration 500 / 1000: loss 1.698016  
iteration 600 / 1000: loss 1.694025  
iteration 700 / 1000: loss 1.521366  
iteration 800 / 1000: loss 1.605131  
iteration 900 / 1000: loss 1.594005  
For parameters: [3.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.435  
iteration 0 / 1000: loss 2.389172  
iteration 100 / 1000: loss 1.943899  
iteration 200 / 1000: loss 1.964574  
iteration 300 / 1000: loss 1.655812  
iteration 400 / 1000: loss 1.687209  
iteration 500 / 1000: loss 1.422541  
iteration 600 / 1000: loss 1.659894  
iteration 700 / 1000: loss 1.809398  
iteration 800 / 1000: loss 1.628091  
iteration 900 / 1000: loss 1.548941  
For parameters: [3.e-02 5.e+01 5.e+01]  
Validation accuracy: 0.447  
iteration 0 / 1000: loss 2.413365  
iteration 100 / 1000: loss 1.866576  
iteration 200 / 1000: loss 1.730257  
iteration 300 / 1000: loss 1.662296  
iteration 400 / 1000: loss 1.563400  
iteration 500 / 1000: loss 1.623761  
iteration 600 / 1000: loss 1.692550  
iteration 700 / 1000: loss 1.285304  
iteration 800 / 1000: loss 1.429589  
iteration 900 / 1000: loss 1.559710  
For parameters: [3.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.481

iteration 0 / 1000: loss 2.369549  
iteration 100 / 1000: loss 1.744381  
iteration 200 / 1000: loss 1.712059  
iteration 300 / 1000: loss 1.589181  
iteration 400 / 1000: loss 1.518529  
iteration 500 / 1000: loss 1.573149  
iteration 600 / 1000: loss 1.398003  
iteration 700 / 1000: loss 1.294511  
iteration 800 / 1000: loss 1.483313  
iteration 900 / 1000: loss 1.467256  
For parameters: [3.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.48  
iteration 0 / 1000: loss 2.433897  
iteration 100 / 1000: loss 1.942873  
iteration 200 / 1000: loss 1.849292  
iteration 300 / 1000: loss 1.610350  
iteration 400 / 1000: loss 1.602557  
iteration 500 / 1000: loss 1.213208  
iteration 600 / 1000: loss 1.624326  
iteration 700 / 1000: loss 1.665147  
iteration 800 / 1000: loss 1.373463  
iteration 900 / 1000: loss 1.390888  
For parameters: [3.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.459  
iteration 0 / 1000: loss 2.549713  
iteration 100 / 1000: loss 1.921282  
iteration 200 / 1000: loss 1.726941  
iteration 300 / 1000: loss 1.819238  
iteration 400 / 1000: loss 1.521763  
iteration 500 / 1000: loss 1.688051  
iteration 600 / 1000: loss 1.847164  
iteration 700 / 1000: loss 1.636834  
iteration 800 / 1000: loss 1.468065  
iteration 900 / 1000: loss 1.574748  
For parameters: [3.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.423  
iteration 0 / 1000: loss 2.419260  
iteration 100 / 1000: loss 1.816838  
iteration 200 / 1000: loss 1.684026  
iteration 300 / 1000: loss 1.623757  
iteration 400 / 1000: loss 1.623373  
iteration 500 / 1000: loss 1.496840  
iteration 600 / 1000: loss 1.574236  
iteration 700 / 1000: loss 1.358700  
iteration 800 / 1000: loss 1.518648  
iteration 900 / 1000: loss 1.544579  
For parameters: [3.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.485  
iteration 0 / 1000: loss 2.389910  
iteration 100 / 1000: loss 1.677383  
iteration 200 / 1000: loss 1.815267  
iteration 300 / 1000: loss 1.723192  
iteration 400 / 1000: loss 1.675641  
iteration 500 / 1000: loss 1.700449  
iteration 600 / 1000: loss 1.626160  
iteration 700 / 1000: loss 1.608528  
iteration 800 / 1000: loss 1.574901



```
iteration 900 / 1000: loss 1.697080
For parameters: [3.e-02 8.e+01 5.e+01]
Validation accuracy: 0.441
iteration 0 / 1000: loss 2.412420
iteration 100 / 1000: loss 1.898045
iteration 200 / 1000: loss 1.734733
iteration 300 / 1000: loss 1.649473
iteration 400 / 1000: loss 1.453475
iteration 500 / 1000: loss 1.713627
iteration 600 / 1000: loss 1.467008
iteration 700 / 1000: loss 1.527426
iteration 800 / 1000: loss 1.672172
iteration 900 / 1000: loss 1.361769
For parameters: [3.e-02 1.e+02 1.e+02]
Validation accuracy: 0.475
iteration 0 / 1000: loss 2.315362
iteration 100 / 1000: loss 1.653302
iteration 200 / 1000: loss 1.606265
iteration 300 / 1000: loss 1.582040
iteration 400 / 1000: loss 1.599255
iteration 500 / 1000: loss 1.589134
iteration 600 / 1000: loss 1.472316
iteration 700 / 1000: loss 1.517007
iteration 800 / 1000: loss 1.286834
iteration 900 / 1000: loss 1.183258
For parameters: [3.0e-02 1.2e+02 2.0e+02]
Validation accuracy: 0.478
The best combination of parameters is: learning rate = 0.1 , hidden layer si
ze = 100 , batch size = 200
For this combination, the validation accuracy is 0.507
```

## Three-layer Sigmoid Activation Network

In [18]: `best_3layer_sigmoid = None` # store the best model into this

```
#####  
###  
# TODO: Tune hyperparameters using the validation set. Store your best trained  
#  
# model in best_3layer_sigmoid.  
#  
#####  
###  
  
#Initialize "best" validation accuracy  
  
best_val_acc = 0  
  
#Initialize vector to store all validation accuracies  
  
val_acc_vec = np.zeros(len(parameters_matrix))  
  
#Cycle through the matrix of parameters  
  
for i in range(len(parameters_matrix)):  
  
    input_size = 32 * 32 * 3  
    num_layers = 3  
    hidden_size = int(parameters_matrix[i,1])  
    hidden_sizes = [hidden_size]*(num_layers-1)  
    num_classes = 10  
    learning_rate = parameters_matrix[i,0]  
    learning_rate_decay = 0.95  
    batch_size = int(parameters_matrix[i,2])  
    net = NeuralNetwork(input_size, hidden_sizes, num_classes, num_layers, nonli  
nearity='sigmoid')  
  
    # Train the network  
    stats = net.train(X_train, y_train, X_val, y_val,  
                      num_iters=1000, batch_size=batch_size,  
                      learning_rate=learning_rate, learning_rate_decay=learning_rate_d  
ecay,  
                      reg=0.0, verbose=True)  
  
    # Predict on the validation set  
    val_acc = (net.predict(X_val) == y_val).mean()  
  
    print('For parameters: ', parameters_matrix[i])  
    print('Validation accuracy: ', val_acc)  
  
    #Store the validation accuracies  
  
    val_acc_vec[i] = val_acc  
  
    #Save the best model and stats  
  
    if val_acc > best_val_acc:  
  
        best_val_acc = val_acc
```

```
best_3layer_sigmoid = net  
best_stats = stats
```

```
#Print the best combination of trial parameters:
```

```
ind = np.argmax(val_acc_vec)
```

```
print('The best combination of parameters is: learning rate = ',parameters_mat  
rix[ind,0], ', hidden layer size = ',int(parameters_matrix[ind,1]), ', batch s  
ize = ',int(parameters_matrix[ind,2]))  
print('For this combination, the validation accuracy is ', val_acc_vec[ind])
```

iteration 0 / 1000: loss 2.341498  
iteration 100 / 1000: loss 2.207852  
iteration 200 / 1000: loss 2.139997  
iteration 300 / 1000: loss 2.051488  
iteration 400 / 1000: loss 2.049709  
iteration 500 / 1000: loss 2.104036  
iteration 600 / 1000: loss 2.050512  
iteration 700 / 1000: loss 1.851410  
iteration 800 / 1000: loss 1.940754  
iteration 900 / 1000: loss 2.018969  
For parameters: [ 0.1 20. 50. ]  
Validation accuracy: 0.357  
iteration 0 / 1000: loss 2.448612  
iteration 100 / 1000: loss 2.077613  
iteration 200 / 1000: loss 2.039883  
iteration 300 / 1000: loss 2.033507  
iteration 400 / 1000: loss 1.895880  
iteration 500 / 1000: loss 1.932328  
iteration 600 / 1000: loss 1.925898  
iteration 700 / 1000: loss 1.838447  
iteration 800 / 1000: loss 1.856493  
iteration 900 / 1000: loss 1.591386  
For parameters: [ 0.1 50. 100. ]  
Validation accuracy: 0.386  
iteration 0 / 1000: loss 2.364265  
iteration 100 / 1000: loss 2.081808  
iteration 200 / 1000: loss 2.016848  
iteration 300 / 1000: loss 1.931935  
iteration 400 / 1000: loss 1.859091  
iteration 500 / 1000: loss 1.844422  
iteration 600 / 1000: loss 1.795662  
iteration 700 / 1000: loss 1.825233  
iteration 800 / 1000: loss 1.681379  
iteration 900 / 1000: loss 1.788236  
For parameters: [1.e-01 8.e+01 2.e+02]  
Validation accuracy: 0.414  
iteration 0 / 1000: loss 2.413968  
iteration 100 / 1000: loss 2.127262  
iteration 200 / 1000: loss 2.028107  
iteration 300 / 1000: loss 1.918618  
iteration 400 / 1000: loss 1.834029  
iteration 500 / 1000: loss 1.918868  
iteration 600 / 1000: loss 1.710111  
iteration 700 / 1000: loss 1.979391  
iteration 800 / 1000: loss 1.606796  
iteration 900 / 1000: loss 1.870395  
For parameters: [ 0.1 100. 50. ]  
Validation accuracy: 0.402  
iteration 0 / 1000: loss 2.389077  
iteration 100 / 1000: loss 2.102345  
iteration 200 / 1000: loss 1.956554  
iteration 300 / 1000: loss 1.985965  
iteration 400 / 1000: loss 1.777410  
iteration 500 / 1000: loss 1.974525  
iteration 600 / 1000: loss 1.737353  
iteration 700 / 1000: loss 1.668831  
iteration 800 / 1000: loss 1.772797

iteration 900 / 1000: loss 1.746701  
For parameters: [1.0e-01 1.2e+02 1.0e+02]  
Validation accuracy: 0.384  
iteration 0 / 1000: loss 2.340397  
iteration 100 / 1000: loss 2.210570  
iteration 200 / 1000: loss 2.149262  
iteration 300 / 1000: loss 2.050630  
iteration 400 / 1000: loss 2.020849  
iteration 500 / 1000: loss 1.944951  
iteration 600 / 1000: loss 1.877223  
iteration 700 / 1000: loss 1.891415  
iteration 800 / 1000: loss 1.842731  
iteration 900 / 1000: loss 1.780380  
For parameters: [1.e-01 2.e+01 2.e+02]  
Validation accuracy: 0.364  
iteration 0 / 1000: loss 2.370122  
iteration 100 / 1000: loss 2.148976  
iteration 200 / 1000: loss 2.010571  
iteration 300 / 1000: loss 2.057767  
iteration 400 / 1000: loss 1.919012  
iteration 500 / 1000: loss 1.860668  
iteration 600 / 1000: loss 1.922120  
iteration 700 / 1000: loss 1.898974  
iteration 800 / 1000: loss 1.886093  
iteration 900 / 1000: loss 1.821627  
For parameters: [ 0.1 50. 50. ]  
Validation accuracy: 0.374  
iteration 0 / 1000: loss 2.479070  
iteration 100 / 1000: loss 2.042379  
iteration 200 / 1000: loss 2.066532  
iteration 300 / 1000: loss 1.969494  
iteration 400 / 1000: loss 1.981710  
iteration 500 / 1000: loss 1.830800  
iteration 600 / 1000: loss 1.884657  
iteration 700 / 1000: loss 1.804319  
iteration 800 / 1000: loss 1.784642  
iteration 900 / 1000: loss 1.819560  
For parameters: [ 0.1 80. 100. ]  
Validation accuracy: 0.379  
iteration 0 / 1000: loss 2.458014  
iteration 100 / 1000: loss 2.117269  
iteration 200 / 1000: loss 2.002774  
iteration 300 / 1000: loss 1.930603  
iteration 400 / 1000: loss 1.880399  
iteration 500 / 1000: loss 1.828451  
iteration 600 / 1000: loss 1.828094  
iteration 700 / 1000: loss 1.844087  
iteration 800 / 1000: loss 1.813307  
iteration 900 / 1000: loss 1.770259  
For parameters: [1.e-01 1.e+02 2.e+02]  
Validation accuracy: 0.371  
iteration 0 / 1000: loss 2.481327  
iteration 100 / 1000: loss 2.065398  
iteration 200 / 1000: loss 1.981070  
iteration 300 / 1000: loss 1.781368  
iteration 400 / 1000: loss 1.890408  
iteration 500 / 1000: loss 1.905495

iteration 600 / 1000: loss 1.839487  
iteration 700 / 1000: loss 1.854714  
iteration 800 / 1000: loss 1.843192  
iteration 900 / 1000: loss 1.760606  
For parameters: [1.0e-01 1.2e+02 5.0e+01]  
Validation accuracy: 0.406  
iteration 0 / 1000: loss 2.421596  
iteration 100 / 1000: loss 2.220491  
iteration 200 / 1000: loss 2.189993  
iteration 300 / 1000: loss 2.026762  
iteration 400 / 1000: loss 2.038135  
iteration 500 / 1000: loss 2.111229  
iteration 600 / 1000: loss 1.932581  
iteration 700 / 1000: loss 1.892606  
iteration 800 / 1000: loss 1.821492  
iteration 900 / 1000: loss 1.926973  
For parameters: [ 0.1 20. 100. ]  
Validation accuracy: 0.355  
iteration 0 / 1000: loss 2.349100  
iteration 100 / 1000: loss 2.099988  
iteration 200 / 1000: loss 2.061398  
iteration 300 / 1000: loss 1.875662  
iteration 400 / 1000: loss 1.934079  
iteration 500 / 1000: loss 1.871254  
iteration 600 / 1000: loss 1.806076  
iteration 700 / 1000: loss 1.941981  
iteration 800 / 1000: loss 1.917081  
iteration 900 / 1000: loss 1.861327  
For parameters: [1.e-01 5.e+01 2.e+02]  
Validation accuracy: 0.382  
iteration 0 / 1000: loss 2.370517  
iteration 100 / 1000: loss 2.115029  
iteration 200 / 1000: loss 2.169580  
iteration 300 / 1000: loss 2.089761  
iteration 400 / 1000: loss 1.956986  
iteration 500 / 1000: loss 1.783870  
iteration 600 / 1000: loss 1.788431  
iteration 700 / 1000: loss 1.675544  
iteration 800 / 1000: loss 1.800275  
iteration 900 / 1000: loss 1.747724  
For parameters: [ 0.1 80. 50. ]  
Validation accuracy: 0.389  
iteration 0 / 1000: loss 2.424740  
iteration 100 / 1000: loss 2.076541  
iteration 200 / 1000: loss 1.857389  
iteration 300 / 1000: loss 1.913411  
iteration 400 / 1000: loss 1.768074  
iteration 500 / 1000: loss 1.808753  
iteration 600 / 1000: loss 1.815828  
iteration 700 / 1000: loss 1.783271  
iteration 800 / 1000: loss 1.657926  
iteration 900 / 1000: loss 1.786941  
For parameters: [ 0.1 100. 100. ]  
Validation accuracy: 0.388  
iteration 0 / 1000: loss 2.414725  
iteration 100 / 1000: loss 2.075634  
iteration 200 / 1000: loss 1.901399

iteration 300 / 1000: loss 1.885077  
iteration 400 / 1000: loss 1.822682  
iteration 500 / 1000: loss 1.898381  
iteration 600 / 1000: loss 1.792772  
iteration 700 / 1000: loss 1.765559  
iteration 800 / 1000: loss 1.855042  
iteration 900 / 1000: loss 1.734327  
For parameters: [1.0e-01 1.2e+02 2.0e+02]  
Validation accuracy: 0.407  
iteration 0 / 1000: loss 2.351040  
iteration 100 / 1000: loss 2.266038  
iteration 200 / 1000: loss 2.194641  
iteration 300 / 1000: loss 2.204253  
iteration 400 / 1000: loss 2.172737  
iteration 500 / 1000: loss 2.078930  
iteration 600 / 1000: loss 2.101879  
iteration 700 / 1000: loss 1.967918  
iteration 800 / 1000: loss 1.987544  
iteration 900 / 1000: loss 1.879330  
For parameters: [ 0.05 20. 50. ]  
Validation accuracy: 0.329  
iteration 0 / 1000: loss 2.370049  
iteration 100 / 1000: loss 2.168831  
iteration 200 / 1000: loss 2.093210  
iteration 300 / 1000: loss 2.053266  
iteration 400 / 1000: loss 2.101964  
iteration 500 / 1000: loss 1.908137  
iteration 600 / 1000: loss 1.999922  
iteration 700 / 1000: loss 1.909557  
iteration 800 / 1000: loss 1.847564  
iteration 900 / 1000: loss 1.904145  
For parameters: [5.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.351  
iteration 0 / 1000: loss 2.412317  
iteration 100 / 1000: loss 2.198026  
iteration 200 / 1000: loss 2.095225  
iteration 300 / 1000: loss 1.976668  
iteration 400 / 1000: loss 1.941050  
iteration 500 / 1000: loss 1.910940  
iteration 600 / 1000: loss 1.878348  
iteration 700 / 1000: loss 1.938718  
iteration 800 / 1000: loss 1.935688  
iteration 900 / 1000: loss 1.807334  
For parameters: [5.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.352  
iteration 0 / 1000: loss 2.486944  
iteration 100 / 1000: loss 2.256717  
iteration 200 / 1000: loss 2.028991  
iteration 300 / 1000: loss 1.982665  
iteration 400 / 1000: loss 2.100781  
iteration 500 / 1000: loss 1.986462  
iteration 600 / 1000: loss 1.942680  
iteration 700 / 1000: loss 1.698867  
iteration 800 / 1000: loss 1.729140  
iteration 900 / 1000: loss 1.840234  
For parameters: [5.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.38

iteration 0 / 1000: loss 2.315558  
iteration 100 / 1000: loss 2.128884  
iteration 200 / 1000: loss 2.012528  
iteration 300 / 1000: loss 2.068834  
iteration 400 / 1000: loss 1.972363  
iteration 500 / 1000: loss 1.995456  
iteration 600 / 1000: loss 1.882271  
iteration 700 / 1000: loss 1.860277  
iteration 800 / 1000: loss 1.784486  
iteration 900 / 1000: loss 1.893442  
For parameters: [5.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.374  
iteration 0 / 1000: loss 2.332475  
iteration 100 / 1000: loss 2.225142  
iteration 200 / 1000: loss 2.162154  
iteration 300 / 1000: loss 2.135651  
iteration 400 / 1000: loss 2.100617  
iteration 500 / 1000: loss 2.048210  
iteration 600 / 1000: loss 2.094557  
iteration 700 / 1000: loss 1.985015  
iteration 800 / 1000: loss 2.026365  
iteration 900 / 1000: loss 1.992307  
For parameters: [5.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.293  
iteration 0 / 1000: loss 2.334120  
iteration 100 / 1000: loss 2.162681  
iteration 200 / 1000: loss 2.091840  
iteration 300 / 1000: loss 2.089505  
iteration 400 / 1000: loss 1.941406  
iteration 500 / 1000: loss 1.974974  
iteration 600 / 1000: loss 1.919954  
iteration 700 / 1000: loss 2.042265  
iteration 800 / 1000: loss 1.930522  
iteration 900 / 1000: loss 1.939204  
For parameters: [ 0.05 50. 50. ]  
Validation accuracy: 0.353  
iteration 0 / 1000: loss 2.384247  
iteration 100 / 1000: loss 2.216204  
iteration 200 / 1000: loss 2.162714  
iteration 300 / 1000: loss 1.987644  
iteration 400 / 1000: loss 2.003775  
iteration 500 / 1000: loss 1.977043  
iteration 600 / 1000: loss 1.948022  
iteration 700 / 1000: loss 1.917032  
iteration 800 / 1000: loss 1.861673  
iteration 900 / 1000: loss 1.855129  
For parameters: [5.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.358  
iteration 0 / 1000: loss 2.456170  
iteration 100 / 1000: loss 2.143607  
iteration 200 / 1000: loss 2.061057  
iteration 300 / 1000: loss 2.011322  
iteration 400 / 1000: loss 1.956127  
iteration 500 / 1000: loss 1.946235  
iteration 600 / 1000: loss 1.872850  
iteration 700 / 1000: loss 1.808731  
iteration 800 / 1000: loss 1.926650



iteration 900 / 1000: loss 1.799540  
For parameters: [5.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.372  
iteration 0 / 1000: loss 2.358424  
iteration 100 / 1000: loss 2.158301  
iteration 200 / 1000: loss 2.082831  
iteration 300 / 1000: loss 1.996793  
iteration 400 / 1000: loss 1.883716  
iteration 500 / 1000: loss 2.097584  
iteration 600 / 1000: loss 1.791184  
iteration 700 / 1000: loss 1.875459  
iteration 800 / 1000: loss 1.991684  
iteration 900 / 1000: loss 1.684003  
For parameters: [5.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.364  
iteration 0 / 1000: loss 2.426748  
iteration 100 / 1000: loss 2.212196  
iteration 200 / 1000: loss 2.153461  
iteration 300 / 1000: loss 2.140821  
iteration 400 / 1000: loss 2.109336  
iteration 500 / 1000: loss 2.098734  
iteration 600 / 1000: loss 2.096055  
iteration 700 / 1000: loss 2.009370  
iteration 800 / 1000: loss 1.954278  
iteration 900 / 1000: loss 2.041460  
For parameters: [5.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.309  
iteration 0 / 1000: loss 2.488710  
iteration 100 / 1000: loss 2.201617  
iteration 200 / 1000: loss 2.121595  
iteration 300 / 1000: loss 2.083578  
iteration 400 / 1000: loss 2.024399  
iteration 500 / 1000: loss 1.957256  
iteration 600 / 1000: loss 1.956924  
iteration 700 / 1000: loss 1.985621  
iteration 800 / 1000: loss 1.922655  
iteration 900 / 1000: loss 1.849824  
For parameters: [5.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.357  
iteration 0 / 1000: loss 2.431763  
iteration 100 / 1000: loss 2.192848  
iteration 200 / 1000: loss 2.067292  
iteration 300 / 1000: loss 2.047826  
iteration 400 / 1000: loss 2.025818  
iteration 500 / 1000: loss 2.066377  
iteration 600 / 1000: loss 1.940036  
iteration 700 / 1000: loss 1.756844  
iteration 800 / 1000: loss 2.036926  
iteration 900 / 1000: loss 2.043059  
For parameters: [5.e-02 8.e+01 5.e+01]  
Validation accuracy: 0.369  
iteration 0 / 1000: loss 2.426608  
iteration 100 / 1000: loss 2.179206  
iteration 200 / 1000: loss 2.135026  
iteration 300 / 1000: loss 2.053648  
iteration 400 / 1000: loss 1.964938  
iteration 500 / 1000: loss 1.827273

iteration 600 / 1000: loss 1.957723  
iteration 700 / 1000: loss 1.860583  
iteration 800 / 1000: loss 1.865474  
iteration 900 / 1000: loss 1.856841  
For parameters: [5.e-02 1.e+02 1.e+02]  
Validation accuracy: 0.34  
iteration 0 / 1000: loss 2.322078  
iteration 100 / 1000: loss 2.165835  
iteration 200 / 1000: loss 2.043643  
iteration 300 / 1000: loss 2.044443  
iteration 400 / 1000: loss 1.978318  
iteration 500 / 1000: loss 1.970712  
iteration 600 / 1000: loss 1.925278  
iteration 700 / 1000: loss 1.862245  
iteration 800 / 1000: loss 1.861564  
iteration 900 / 1000: loss 1.921277  
For parameters: [5.0e-02 1.2e+02 2.0e+02]  
Validation accuracy: 0.37  
iteration 0 / 1000: loss 2.410755  
iteration 100 / 1000: loss 2.254633  
iteration 200 / 1000: loss 2.202615  
iteration 300 / 1000: loss 2.233503  
iteration 400 / 1000: loss 2.187810  
iteration 500 / 1000: loss 2.161803  
iteration 600 / 1000: loss 2.165336  
iteration 700 / 1000: loss 2.112059  
iteration 800 / 1000: loss 2.169692  
iteration 900 / 1000: loss 2.045631  
For parameters: [3.e-02 2.e+01 5.e+01]  
Validation accuracy: 0.296  
iteration 0 / 1000: loss 2.517343  
iteration 100 / 1000: loss 2.243674  
iteration 200 / 1000: loss 2.225400  
iteration 300 / 1000: loss 2.159400  
iteration 400 / 1000: loss 2.138322  
iteration 500 / 1000: loss 2.140139  
iteration 600 / 1000: loss 2.076396  
iteration 700 / 1000: loss 2.086987  
iteration 800 / 1000: loss 2.005066  
iteration 900 / 1000: loss 1.967004  
For parameters: [3.e-02 5.e+01 1.e+02]  
Validation accuracy: 0.327  
iteration 0 / 1000: loss 2.584331  
iteration 100 / 1000: loss 2.214143  
iteration 200 / 1000: loss 2.159015  
iteration 300 / 1000: loss 2.138768  
iteration 400 / 1000: loss 2.062011  
iteration 500 / 1000: loss 1.975079  
iteration 600 / 1000: loss 1.978128  
iteration 700 / 1000: loss 2.004669  
iteration 800 / 1000: loss 1.929064  
iteration 900 / 1000: loss 1.901959  
For parameters: [3.e-02 8.e+01 2.e+02]  
Validation accuracy: 0.324  
iteration 0 / 1000: loss 2.355964  
iteration 100 / 1000: loss 2.176187  
iteration 200 / 1000: loss 2.122950

iteration 300 / 1000: loss 2.063921  
iteration 400 / 1000: loss 1.967118  
iteration 500 / 1000: loss 2.067608  
iteration 600 / 1000: loss 1.831401  
iteration 700 / 1000: loss 1.958748  
iteration 800 / 1000: loss 1.961925  
iteration 900 / 1000: loss 1.914812  
For parameters: [3.e-02 1.e+02 5.e+01]  
Validation accuracy: 0.342  
iteration 0 / 1000: loss 2.432296  
iteration 100 / 1000: loss 2.212569  
iteration 200 / 1000: loss 2.073832  
iteration 300 / 1000: loss 2.054449  
iteration 400 / 1000: loss 2.021154  
iteration 500 / 1000: loss 1.895779  
iteration 600 / 1000: loss 2.001518  
iteration 700 / 1000: loss 2.004096  
iteration 800 / 1000: loss 1.989022  
iteration 900 / 1000: loss 1.958868  
For parameters: [3.0e-02 1.2e+02 1.0e+02]  
Validation accuracy: 0.347  
iteration 0 / 1000: loss 2.410058  
iteration 100 / 1000: loss 2.245920  
iteration 200 / 1000: loss 2.223053  
iteration 300 / 1000: loss 2.203342  
iteration 400 / 1000: loss 2.159479  
iteration 500 / 1000: loss 2.133978  
iteration 600 / 1000: loss 2.130835  
iteration 700 / 1000: loss 2.098385  
iteration 800 / 1000: loss 2.140647  
iteration 900 / 1000: loss 2.082727  
For parameters: [3.e-02 2.e+01 2.e+02]  
Validation accuracy: 0.274  
iteration 0 / 1000: loss 2.278432  
iteration 100 / 1000: loss 2.213406  
iteration 200 / 1000: loss 2.208855  
iteration 300 / 1000: loss 2.066700  
iteration 400 / 1000: loss 2.051404  
iteration 500 / 1000: loss 2.077448  
iteration 600 / 1000: loss 2.030642  
iteration 700 / 1000: loss 1.890243  
iteration 800 / 1000: loss 2.013143  
iteration 900 / 1000: loss 1.944813  
For parameters: [3.e-02 5.e+01 5.e+01]  
Validation accuracy: 0.323  
iteration 0 / 1000: loss 2.359923  
iteration 100 / 1000: loss 2.203286  
iteration 200 / 1000: loss 2.120148  
iteration 300 / 1000: loss 2.125394  
iteration 400 / 1000: loss 2.053189  
iteration 500 / 1000: loss 1.995812  
iteration 600 / 1000: loss 2.018539  
iteration 700 / 1000: loss 2.050029  
iteration 800 / 1000: loss 1.940142  
iteration 900 / 1000: loss 1.933720  
For parameters: [3.e-02 8.e+01 1.e+02]  
Validation accuracy: 0.334

iteration 0 / 1000: loss 2.405881  
iteration 100 / 1000: loss 2.213258  
iteration 200 / 1000: loss 2.137062  
iteration 300 / 1000: loss 2.156366  
iteration 400 / 1000: loss 2.055444  
iteration 500 / 1000: loss 1.980912  
iteration 600 / 1000: loss 1.989552  
iteration 700 / 1000: loss 2.006236  
iteration 800 / 1000: loss 1.967418  
iteration 900 / 1000: loss 1.890400  
For parameters: [3.e-02 1.e+02 2.e+02]  
Validation accuracy: 0.342  
iteration 0 / 1000: loss 2.374897  
iteration 100 / 1000: loss 2.138628  
iteration 200 / 1000: loss 2.068674  
iteration 300 / 1000: loss 2.125135  
iteration 400 / 1000: loss 2.040380  
iteration 500 / 1000: loss 1.963412  
iteration 600 / 1000: loss 1.904206  
iteration 700 / 1000: loss 1.920922  
iteration 800 / 1000: loss 1.934204  
iteration 900 / 1000: loss 1.865893  
For parameters: [3.0e-02 1.2e+02 5.0e+01]  
Validation accuracy: 0.355  
iteration 0 / 1000: loss 2.496609  
iteration 100 / 1000: loss 2.246303  
iteration 200 / 1000: loss 2.158703  
iteration 300 / 1000: loss 2.165596  
iteration 400 / 1000: loss 2.142796  
iteration 500 / 1000: loss 2.118159  
iteration 600 / 1000: loss 2.097155  
iteration 700 / 1000: loss 2.030318  
iteration 800 / 1000: loss 2.074157  
iteration 900 / 1000: loss 2.019326  
For parameters: [3.e-02 2.e+01 1.e+02]  
Validation accuracy: 0.306  
iteration 0 / 1000: loss 2.427046  
iteration 100 / 1000: loss 2.260295  
iteration 200 / 1000: loss 2.188190  
iteration 300 / 1000: loss 2.178611  
iteration 400 / 1000: loss 2.136770  
iteration 500 / 1000: loss 2.090595  
iteration 600 / 1000: loss 2.073835  
iteration 700 / 1000: loss 2.018430  
iteration 800 / 1000: loss 2.010283  
iteration 900 / 1000: loss 2.006048  
For parameters: [3.e-02 5.e+01 2.e+02]  
Validation accuracy: 0.327  
iteration 0 / 1000: loss 2.552201  
iteration 100 / 1000: loss 2.187638  
iteration 200 / 1000: loss 2.124987  
iteration 300 / 1000: loss 2.150177  
iteration 400 / 1000: loss 2.011614  
iteration 500 / 1000: loss 2.083815  
iteration 600 / 1000: loss 1.995669  
iteration 700 / 1000: loss 2.041478  
iteration 800 / 1000: loss 1.878195

```

iteration 900 / 1000: loss 1.901249
For parameters: [3.e-02 8.e+01 5.e+01]
Validation accuracy: 0.346
iteration 0 / 1000: loss 2.352675
iteration 100 / 1000: loss 2.162777
iteration 200 / 1000: loss 2.103423
iteration 300 / 1000: loss 2.114047
iteration 400 / 1000: loss 2.057731
iteration 500 / 1000: loss 1.995899
iteration 600 / 1000: loss 2.036591
iteration 700 / 1000: loss 1.996298
iteration 800 / 1000: loss 1.986193
iteration 900 / 1000: loss 1.962722
For parameters: [3.e-02 1.e+02 1.e+02]
Validation accuracy: 0.354
iteration 0 / 1000: loss 2.339715
iteration 100 / 1000: loss 2.198587
iteration 200 / 1000: loss 2.154298
iteration 300 / 1000: loss 2.091658
iteration 400 / 1000: loss 2.037219
iteration 500 / 1000: loss 2.012031
iteration 600 / 1000: loss 1.972614
iteration 700 / 1000: loss 1.979341
iteration 800 / 1000: loss 1.894481
iteration 900 / 1000: loss 1.896944
For parameters: [3.0e-02 1.2e+02 2.0e+02]
Validation accuracy: 0.334
The best combination of parameters is: learning rate = 0.1 , hidden layer si
ze = 80 , batch size = 200
For this combination, the validation accuracy is 0.414

```

## Run on the test set

When you are done experimenting, you should evaluate your final trained networks on the test set.

```

In [10]: print('Two-layer relu')
test_acc = (best_2layer_relu.predict(X_test) == y_test).mean()
print('Test accuracy: ', test_acc)

```

```

Two-layer relu
Test accuracy: 0.5038

```

```

In [14]: print('Two-layer sigmoid')
test_acc = (best_2layer_sigmoid.predict(X_test) == y_test).mean()
print('Test accuracy: ', test_acc)

```

```

Two-layer sigmoid
Test accuracy: 0.4354

```

```
In [17]: print('Three-layer relu')
test_acc = (best_3layer_relu.predict(X_test) == y_test).mean()
print('Test accuracy: ', test_acc)
```

Three-layer relu  
Test accuracy: 0.505

```
In [19]: print('Three-layer sigmoid')
test_acc = (best_3layer_sigmoid.predict(X_test) == y_test).mean()
print('Test accuracy: ', test_acc)
```

Three-layer sigmoid  
Test accuracy: 0.3964

## Kaggle output

Once you are satisfied with your solution and test accuracy output a file to submit your test set predictions to the Kaggle for Assignment 2 Neural Network. Use the following code to do so:

```
In [12]: output_submission_csv('nn_2layer_relu_submission.csv', best_2layer_relu.predict(X_test))
```

```
In [15]: output_submission_csv('nn_2layer_sigmoid_submission.csv', best_2layer_sigmoid.predict(X_test))
```

```
In [17]: output_submission_csv('nn_3layer_relu_submission.csv', best_3layer_relu.predict(X_test))
```

```
In [18]: output_submission_csv('nn_3layer_sigmoid_submission.csv', best_3layer_sigmoid.predict(X_test))
```